



Name : Tavhare Ruchita Sharad

class : BE AI&DS

Roll No.: 61

Subject: Computer Laboratory-I(ML)

Title: Predict the price of the Uber ride from a given pickup point to the agreed drop-off location. Perform following tasks:

1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and ridge, Lasso regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [3]: # 1. Load and Pre-process the Dataset
df = pd.read_csv('uber.csv') # Change the name to the correct csv file
```

```
In [6]: print("Initial Data Info:")
print(df.info())
print(df.head())
```

```

Initial Data Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Unnamed: 0         200000 non-null   int64  
 1   key               200000 non-null   object  
 2   fare_amount       200000 non-null   float64 
 3   pickup_datetime   200000 non-null   object  
 4   pickup_longitude  200000 non-null   float64 
 5   pickup_latitude   200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude  199999 non-null   float64 
 8   passenger_count   200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
None
      Unnamed: 0           key  fare_amount \
0    24238194  2015-05-07 19:52:06.0000003    7.5
1    27835199  2009-07-17 20:04:56.0000002    7.7
2    44984355  2009-08-24 21:45:00.0000001   12.9
3    25894730  2009-06-26 08:22:21.0000001    5.3
4    17610152  2014-08-28 17:47:00.000000188   16.0

      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC          -73.999817        40.738354
1  2009-07-17 20:04:56 UTC          -73.994355        40.728225
2  2009-08-24 21:45:00 UTC          -74.005043        40.740770
3  2009-06-26 08:22:21 UTC          -73.976124        40.790844
4  2014-08-28 17:47:00 UTC          -73.925023        40.744085

      dropoff_longitude  dropoff_latitude  passenger_count
0          -73.999512        40.723217             1
1          -73.994710        40.750325             1
2          -73.962565        40.772647             1
3          -73.965316        40.803349             3
4          -73.973082        40.761247             5

```

```
In [9]: # Drop rows with missing values
df = df.dropna()
```

```
In [11]: # Convert datetime columns
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

```
In [12]: # Feature Engineering: Extract hour, day, month
df['hour'] = df['pickup_datetime'].dt.hour
df['day'] = df['pickup_datetime'].dt.day
df['month'] = df['pickup_datetime'].dt.month
```

```
In [13]: # Remove unneeded columns
df = df.drop(columns=['key', 'pickup_datetime'])
```


```
In [17]: # Calculate distance (Haversine formula)
def haversine(lat1, lon1, lat2, lon2):
    R = 6371 # Earth radius in km
    phi1, phi2 = np.radians(lat1), np.radians(lat2)
    dphi = np.radians(lat2 - lat1)
    dlambd = np.radians(lon2 - lon1)
    a = np.sin(dphi/2)**2 + np.cos(phi1)*np.cos(phi2)*np.sin(dlambd/2)**2
    return 2 * R * np.arcsin(np.sqrt(a))

df['distance_km'] = haversine(df['pickup_latitude'], df['pickup_longitude'],
                               df['dropoff_latitude'], df['dropoff_longitude'])
```

```
In [19]: # Drop original lat/lon columns
df = df.drop(columns=['pickup_latitude', 'pickup_longitude', 'dropoff_latitude',
                      'dropoff_longitude'])
```

```
In [21]: # 2. Identify Outliers (using z-score on fare_amount and distance_km)
from scipy.stats import zscore
df = df[(np.abs(zscore(df[['fare_amount', 'distance_km']]))) < 3).all(axis=1)]
```

```
In [23]: # 3. Check Correlation
plt.figure(figsize=(8, 5))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()
```



```
In [24]: # 4. Regression Models
X = df.drop(columns=['fare_amount'])
y = df['fare_amount']

In [27]: # One-hot encode categorical columns (if any)
X = pd.get_dummies(X, drop_first=True)

In [29]: # Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [31]: # Linear Regression
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)

In [33]: # Ridge Regression
ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)
y_pred_ridge = ridge.predict(X_test)

In [35]: # Lasso Regression
lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
y_pred_lasso = lasso.predict(X_test)

In [37]: # 5. Model Evaluation
def print_scores(model_name, y_true, y_pred):
    print(f"\n{model_name}")
    print(f"R2 Score: {r2_score(y_true, y_pred):.4f}")
    print(f"RMSE: {np.sqrt(mean_squared_error(y_true, y_pred)):.4f}")

In [39]: print_scores("Linear Regression", y_test, y_pred_lr)
print_scores("Ridge Regression", y_test, y_pred_ridge)
print_scores("Lasso Regression", y_test, y_pred_lasso)

Linear Regression
R2 Score: 0.1917
RMSE: 5.8158

Ridge Regression
R2 Score: 0.1917
RMSE: 5.8158

Lasso Regression
R2 Score: 0.1909
RMSE: 5.8187

In [41]: # Optional: Compare visually
plt.figure(figsize=(8,5))
plt.plot(y_test.values, label='True')
plt.plot(y_pred_lr, label='Linear Regression Predicted')
plt.plot(y_pred_ridge, label='Ridge Predicted')
```

```
plt.plot(y_pred_lasso, label='Lasso Predicted')
plt.title('Model Predictions vs True Values (Sample)')
plt.xlabel('Sample')
plt.ylabel('Fare Amount')
plt.legend()
plt.show()
```

