

# CL1-04

July 24, 2025

```
[ ]: """
NAME: Aher Swami Sandip
ROLL NO. 01
COURSE: AI&DS
CLASS: BE
SUB: Computer Laboratory-I (Machine Learning)
"""

[ ]: """
PRACTICAL NO-04:
    Implement K-Means clustering on Iris.csv dataset. Determine the number of clusters
    using the elbow method.
Dataset Link: https://www.kaggle.com/datasets/uciml/iris
"""

[1]: # Import Required Libraries.

[75]: import os
os.environ["OMP_NUM_THREADS"] = "1"

[77]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
import seaborn as sns

[79]: # Load the Iris Dataset.

[81]: df = pd.read_csv("Iris (1).csv")

[83]: df

[83]:   Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm \
0      1          5.1         3.5        1.4         0.2
1      2          4.9         3.0        1.4         0.2
2      3          4.7         3.2        1.3         0.2
3      4          4.6         3.1        1.5         0.2
```

```
4      5      5.0      3.6      1.4      0.2
..    ...
145  146      6.7      3.0      5.2      2.3
146  147      6.3      2.5      5.0      1.9
147  148      6.5      3.0      5.2      2.0
148  149      6.2      3.4      5.4      2.3
149  150      5.9      3.0      5.1      1.8
```

```
Species
0     Iris-setosa
1     Iris-setosa
2     Iris-setosa
3     Iris-setosa
4     Iris-setosa
..
145   ...
145   Iris-virginica
146   Iris-virginica
147   Iris-virginica
148   Iris-virginica
149   Iris-virginica
```

```
[150 rows x 6 columns]
```

```
[85]: df.head()
```

```
[85]:   Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm    Species
0   1           5.1          3.5          1.4          0.2  Iris-setosa
1   2           4.9          3.0          1.4          0.2  Iris-setosa
2   3           4.7          3.2          1.3          0.2  Iris-setosa
3   4           4.6          3.1          1.5          0.2  Iris-setosa
4   5           5.0          3.6          1.4          0.2  Iris-setosa
```

```
[87]: df.tail()
```

```
[87]:   Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm \
145  146           6.7          3.0          5.2          2.3
146  147           6.3          2.5          5.0          1.9
147  148           6.5          3.0          5.2          2.0
148  149           6.2          3.4          5.4          2.3
149  150           5.9          3.0          5.1          1.8
```

```
Species
145  Iris-virginica
146  Iris-virginica
147  Iris-virginica
148  Iris-virginica
149  Iris-virginica
```

```
[89]: len(df)
[89]: 150
[91]: df.shape
[91]: (150, 6)
[93]: df.columns
[93]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',
   'Species'],
   dtype='object')

[95]: for i, col in enumerate(df.columns):
        print(f'column number {1+i} is {col}')
column number 1 is Id
column number 2 is SepalLengthCm
column number 3 is SepalWidthCm
column number 4 is PetalLengthCm
column number 5 is PetalWidthCm
column number 6 is Species

[97]: df.dtypes
[97]: Id          int64
SepalLengthCm    float64
SepalWidthCm     float64
PetalLengthCm    float64
PetalWidthCm     float64
Species         object
dtype: object

[99]: df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Id          150 non-null    int64  
 1   SepalLengthCm 150 non-null    float64 
 2   SepalWidthCm  150 non-null    float64 
 3   PetalLengthCm 150 non-null    float64 
 4   PetalWidthCm  150 non-null    float64 
 5   Species      150 non-null    object  
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

```
[101]: df.describe()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000

```
[103]: # Checking the Missing Values.
```

```
[105]: df.isnull().sum()
```

```
[105]: Id          0  
SepalLengthCm  0  
SepalWidthCm   0  
PetalLengthCm  0  
PetalWidthCm   0  
Species        0  
dtype: int64
```

```
[107]: df.drop('Id', axis = 1, inplace = True)
```

```
[109]: df.head()
```

```
[109]: SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm      Species  
0            5.1          3.5          1.4          0.2  Iris-setosa  
1            4.9          3.0          1.4          0.2  Iris-setosa  
2            4.7          3.2          1.3          0.2  Iris-setosa  
3            4.6          3.1          1.5          0.2  Iris-setosa  
4            5.0          3.6          1.4          0.2  Iris-setosa
```

```
[111]: # K - Means Clustering.
```

```
[113]: df.isnull().sum()
```

```
[113]: SepalLengthCm  0  
SepalWidthCm       0  
PetalLengthCm     0  
PetalWidthCm      0  
Species           0  
dtype: int64
```

```
[115]: df['Species'].value_counts()
```

```
[115]: Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
Name: count, dtype: int64
```

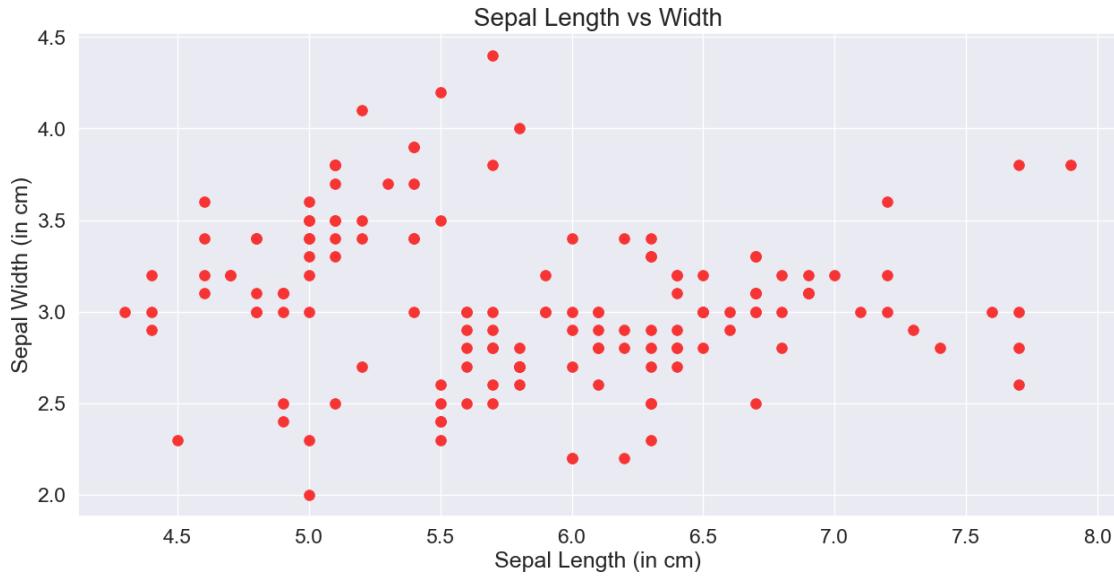
```
[117]: # Splitting into Training and Target Data
# Target Data
target_data = df.iloc[:, 4]
target_data.head()
```

```
[117]: 0    Iris-setosa
1    Iris-setosa
2    Iris-setosa
3    Iris-setosa
4    Iris-setosa
Name: Species, dtype: object
```

```
[119]: # Training Data
clustering_data = df.iloc[:, [0,1,2,3]]
clustering_data.head()
```

```
[119]: SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm
0            5.1          3.5           1.4          0.2
1            4.9          3.0           1.4          0.2
2            4.7          3.2           1.3          0.2
3            4.6          3.1           1.5          0.2
4            5.0          3.6           1.4          0.2
```

```
[121]: fig, ax = plt.subplots(figsize=(15, 7))
sns.set(font_scale=1.5)
ax = sns.scatterplot(x=df['SepalLengthCm'], y=df['SepalWidthCm'], s=70, color="#f73434", edgecolor="#f73434")
ax.set_ylabel('Sepal Width (in cm)')
ax.set_xlabel('Sepal Length (in cm)')
plt.title('Sepal Length vs Width', fontsize=20)
plt.show()
```



```
[123]: # The Elbow Method.
```

```
[125]: import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
from sklearn.cluster import KMeans
wcss = []
for i in range(1, 11):
    km = KMeans(i)
    km.fit(clustering_data)
    wcss.append(km.inertia_)
np.array(wcss)

D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.
    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.
    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.
    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
```

KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.

```

    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.

    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.

    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.

    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.

    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.

    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.

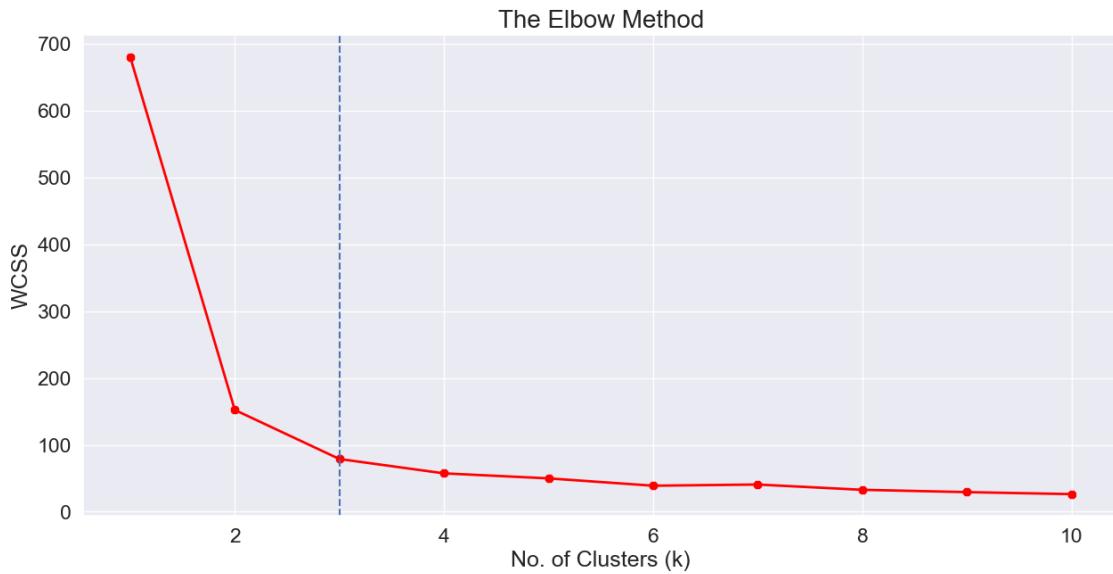
    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less
chunks than available threads. You can avoid it by setting the environment
variable OMP_NUM_THREADS=1.

    warnings.warn(

```

[125]: array([680.8244 , 152.36870648, 78.94084143, 57.43531071, 49.91714056, 38.93873974, 40.7202923 , 32.79068756, 29.33829297, 26.32961079])

[127]: fig, ax = plt.subplots(figsize = (15, 7))
ax = plt.plot(range(1, 11), wcss, linewidth = 2, color = "red", marker = "8")
plt.axvline(x = 3, ls = '--')
plt.ylabel('WCSS')
plt.xlabel('No. of Clusters (k)')
plt.title('The Elbow Method', fontsize = 20)
plt.show()



```
[129]: # Clusters.
```

```
[131]: kms = KMeans(n_clusters = 3, init = 'k-means++')
kms.fit(clustering_data)
KMeans(n_clusters = 3)
clusters = clustering_data.copy()
clusters['Cluster_Prediction'] = kms.fit_predict(clustering_data)
clusters.head()
```

D:\Anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1429: UserWarning:  
KMeans is known to have a memory leak on Windows with MKL, when there are less  
chunks than available threads. You can avoid it by setting the environment  
variable OMP\_NUM\_THREADS=1.

```
    warnings.warn(
D:\Anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning:  
KMeans is known to have a memory leak on Windows with MKL, when there are less  
chunks than available threads. You can avoid it by setting the environment  
variable OMP_NUM_THREADS=1.
    warnings.warn(
```

```
[131]: SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  \
0           5.1          3.5         1.4          0.2
1           4.9          3.0         1.4          0.2
2           4.7          3.2         1.3          0.2
3           4.6          3.1         1.5          0.2
4           5.0          3.6         1.4          0.2
```

Cluster\_Prediction

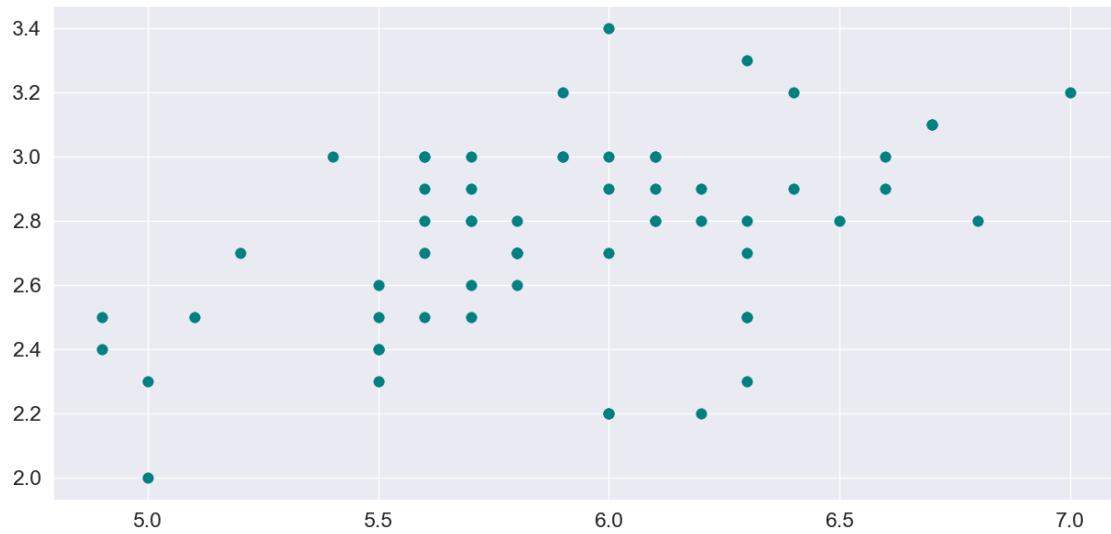
```
0          1  
1          1  
2          1  
3          1  
4          1
```

```
[133]: kms.cluster_centers_
```

```
[133]: array([[5.9016129 , 2.7483871 , 4.39354839, 1.43387097],  
           [5.006      , 3.418      , 1.464      , 0.244      ],  
           [6.85       , 3.07368421, 5.74210526, 2.07105263]])
```

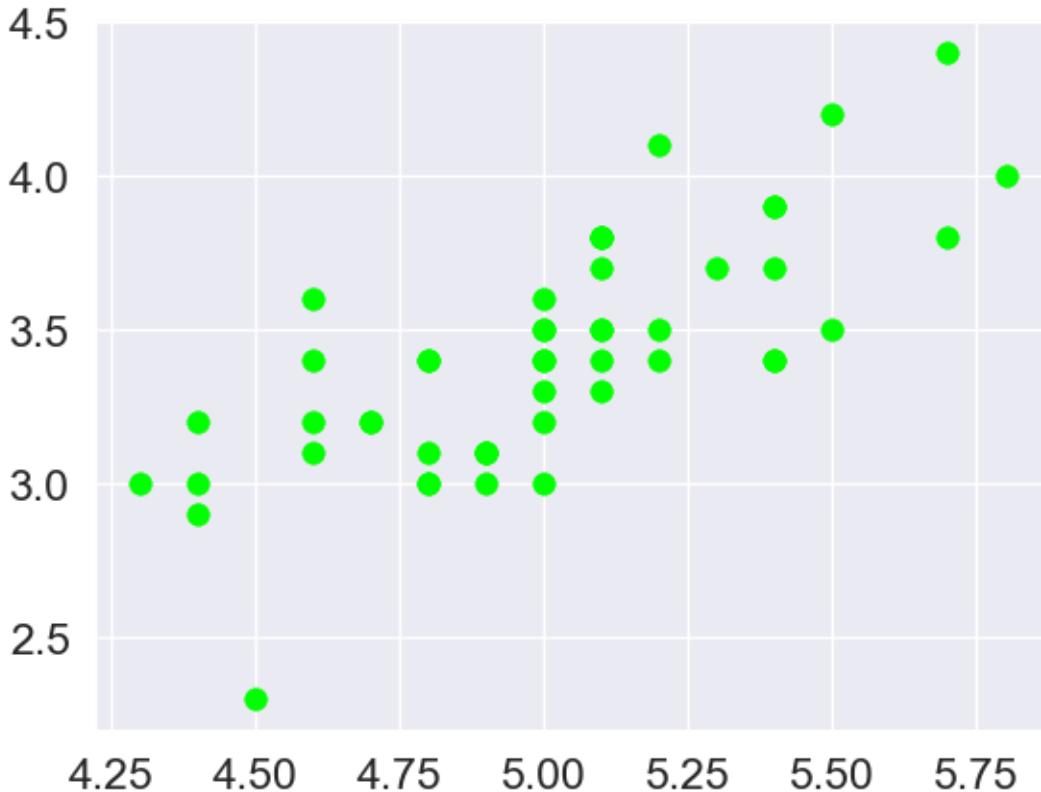
```
[135]: fig, ax = plt.subplots(figsize = (15,7))  
plt.scatter(x = clusters[clusters['Cluster_Prediction'] == 0]['SepalLengthCm'],  
            y = clusters[clusters['Cluster_Prediction'] == 0]['SepalWidthCm'],  
            s = 70, edgecolor = 'teal', linewidth = 0.3, c = 'teal', label =  
            'Iris-versicolor')
```

```
[135]: <matplotlib.collections.PathCollection at 0x1603aeac920>
```



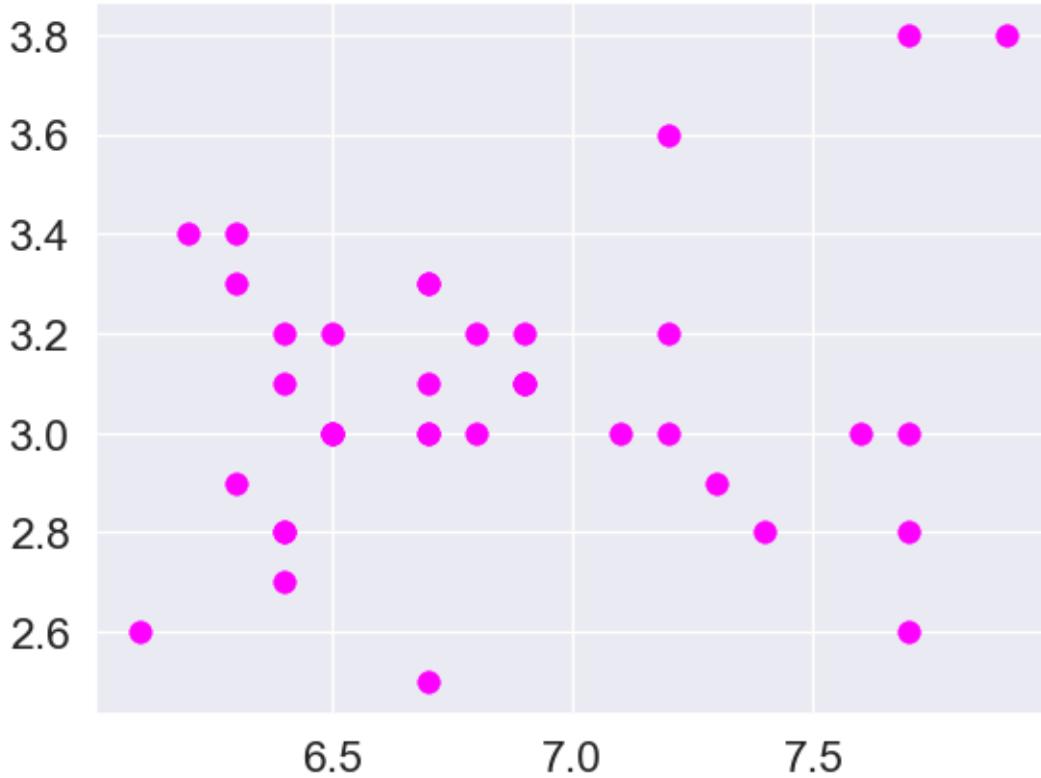
```
[137]: plt.scatter(x = clusters[clusters['Cluster_Prediction'] == 1]['SepalLengthCm'],  
                  y = clusters[clusters['Cluster_Prediction'] == 1]['SepalWidthCm'],  
                  s = 70, edgecolor = 'lime', linewidth = 0.3, c = 'lime', label =  
                  'Iris-setosa')
```

```
[137]: <matplotlib.collections.PathCollection at 0x1603aec5dc0>
```

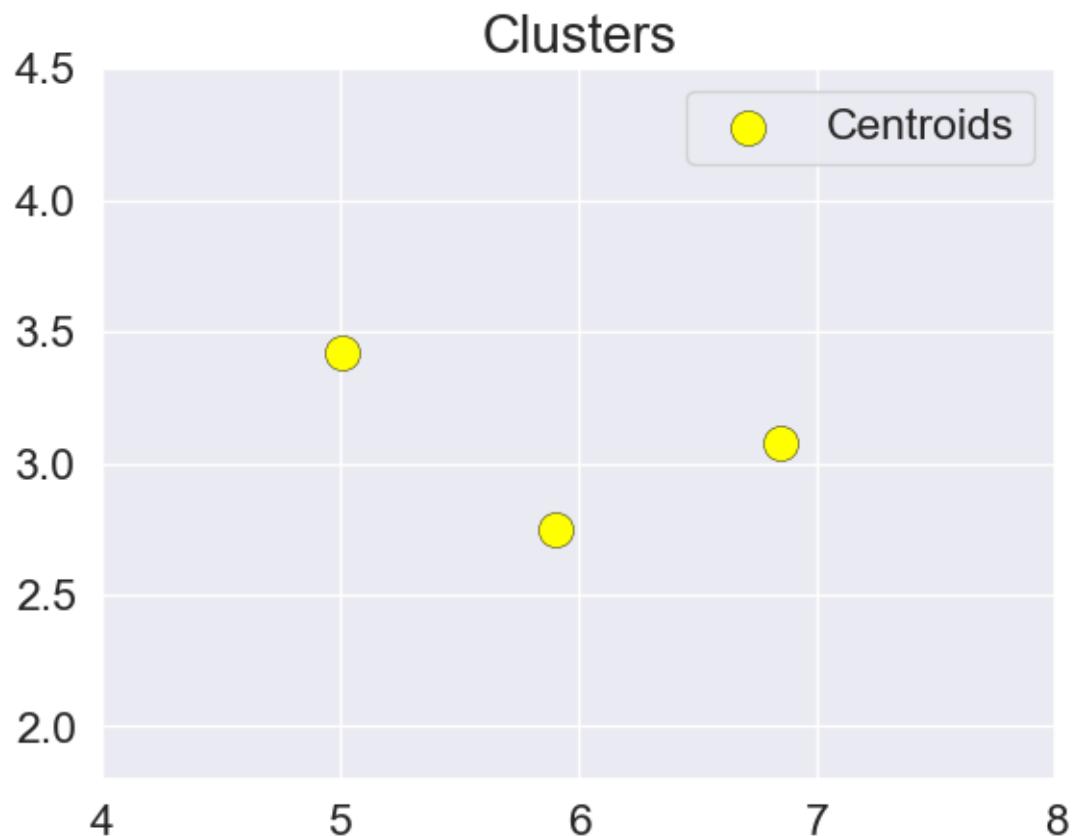


```
[139]: plt.scatter(x = clusters[clusters['Cluster_Prediction'] == 2]['SepalLengthCm'],
                  y = clusters[clusters['Cluster_Prediction'] == 2]['SepalWidthCm'],
                  s = 70, edgecolor = 'magenta', linewidth = 0.3, c = 'magenta',
                  label = 'Iris-virginica')
```

```
[139]: <matplotlib.collections.PathCollection at 0x1603af2ab70>
```



```
[141]: plt.scatter(x = kms.cluster_centers_[:, 0], y = kms.cluster_centers_[:, 1], s = 170, c = 'yellow',
                   label = 'Centroids', edgecolor = 'black', linewidth = 0.3)
plt.legend(loc = 'upper right')
plt.xlim(4, 8)
plt.ylim(1.8, 4.5)
ax.set_ylabel('Sepal Width (in cm)')
ax.set_xlabel('Sepal Length(in cm)')
plt.title('Clusters', fontsize = 20)
plt.show()
```



[ ]: