

CL-I DMV 9

July 16, 2025

```
[1]: ''' NAME:Aher Swami Sandip  
ROLL NO.01  
COURSE: AI&DS  
CLASS: BE  
SUB:Computer Laboratory-I (DMV) '''
```

0.1 Data Cleaning and Preparation

```
[1]: # Problem Statement: Analyzing Customer Churn in a Telecommunications Company  
# Dataset: "Telecom_Customer_Churn.csv"  
# Description: The dataset contains information about customers of a  
# telecommunications  
# company and whether they have churned (i.e., discontinued their services).  
# The dataset  
# includes various attributes of the customers, such as their demographics,  
# usage patterns, and  
# account information. The goal is to perform data cleaning and preparation to  
# gain insights  
# into the factors that contribute to customer churn.  
# Tasks to Perform:
```

0.1.1 1. Import the “Telecom_Customer_Churn.csv” dataset.

```
[1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
df=pd.read_csv("Telco-Customer-Churn.csv")
```

0.1.2 2. Explore the dataset to understand its structure and content.

```
[3]: df.head()
```

```
[3]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService \
0    7590-VHVEG  Female           0      Yes        No         1        No
1    5575-GNVDE    Male           0       No        No        34      Yes
2    3668-QPYBK    Male           0       No        No         2      Yes
3    7795-CFOCW    Male           0       No        No        45        No
```

```

4 9237-HQITU Female          0      No      No      2      Yes
                                         MultipleLines InternetService OnlineSecurity ... DeviceProtection \
0  No phone service                 DSL           No ...           No
1           No                   DSL           Yes ...           Yes
2           No                   DSL           Yes ...           No
3  No phone service                 DSL           Yes ...           Yes
4           No     Fiber optic           No ...           No

TechSupport StreamingTV StreamingMovies          Contract PaperlessBilling \
0           No           No           No Month-to-month           Yes
1           No           No           No One year           No
2           No           No           No Month-to-month           Yes
3           Yes          No           No One year           No
4           No           No           No Month-to-month           Yes

PaymentMethod MonthlyCharges  TotalCharges Churn
0  Electronic check        29.85       29.85      No
1  Mailed check           56.95     1889.5      No
2  Mailed check           53.85      108.15     Yes
3  Bank transfer (automatic)  42.30    1840.75      No
4  Electronic check        70.70      151.65     Yes

```

[5 rows x 21 columns]

[5]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   customerID        7043 non-null   object 
 1   gender             7043 non-null   object 
 2   SeniorCitizen     7043 non-null   int64  
 3   Partner            7043 non-null   object 
 4   Dependents         7043 non-null   object 
 5   tenure             7043 non-null   int64  
 6   PhoneService       7043 non-null   object 
 7   MultipleLines      7043 non-null   object 
 8   InternetService    7043 non-null   object 
 9   OnlineSecurity     7043 non-null   object 
 10  OnlineBackup        7043 non-null   object 
 11  DeviceProtection   7043 non-null   object 
 12  TechSupport         7043 non-null   object 
 13  StreamingTV         7043 non-null   object 
 14  StreamingMovies     7043 non-null   object 
 15  Contract            7043 non-null   object 

```

```
16 PaperlessBilling    7043 non-null   object
17 PaymentMethod       7043 non-null   object
18 MonthlyCharges     7043 non-null   float64
19 TotalCharges        7043 non-null   object
20 Churn               7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
[7]: df.describe()
```

```
[7]:      SeniorCitizen      tenure  MonthlyCharges
count    7043.000000  7043.000000    7043.000000
mean      0.162147   32.371149    64.761692
std       0.368612   24.559481    30.090047
min       0.000000   0.000000    18.250000
25%      0.000000   9.000000    35.500000
50%      0.000000  29.000000    70.350000
75%      0.000000  55.000000   89.850000
max      1.000000  72.000000   118.750000
```

0.1.3 3. Handle missing values in the dataset, deciding on an appropriate strategy.

```
[9]: # Convert "TotalCharges" to numeric, handling any conversion errors.
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')

# Drop rows with missing values in "TotalCharges."
df.dropna(subset=['TotalCharges'], inplace=True)
```

0.1.4 4. Remove any duplicate records from the dataset.

```
[11]: # Remove duplicates
df.drop_duplicates(inplace=True)
```

0.1.5 5. Check for inconsistent data, such as inconsistent formatting or spelling variations, and standardize it.

```
[13]: # Check for inconsistent data by looking at unique values in each column
for column in df.columns:
    print(f"Unique values in {column}:")
    print(df[column].unique())
    print("\n")
```

```
Unique values in customerID:
['7590-VHVEG' '5575-GNVDE' '3668-QPYBK' ... '4801-JZAZL' '8361-LTMKD'
 '3186-AJIEK']
```

```
Unique values in gender:
```

```
['Female' 'Male']
```

```
Unique values in SeniorCitizen:  
[0 1]
```

```
Unique values in Partner:  
['Yes' 'No']
```

```
Unique values in Dependents:  
['No' 'Yes']
```

```
Unique values in tenure:  
[ 1 34  2 45  8 22 10 28 62 13 16 58 49 25 69 52 71 21 12 30 47 72 17 27  
 5 46 11 70 63 43 15 60 18 66  9  3 31 50 64 56  7 42 35 48 29 65 38 68  
 32 55 37 36 41  6  4 33 67 23 57 61 14 20 53 40 59 24 44 19 54 51 26 39]
```

```
Unique values in PhoneService:  
['No' 'Yes']
```

```
Unique values in MultipleLines:  
['No phone service' 'No' 'Yes']
```

```
Unique values in InternetService:  
['DSL' 'Fiber optic' 'No']
```

```
Unique values in OnlineSecurity:  
['No' 'Yes' 'No internet service']
```

```
Unique values in OnlineBackup:  
['Yes' 'No' 'No internet service']
```

```
Unique values in DeviceProtection:  
['No' 'Yes' 'No internet service']
```

```
Unique values in TechSupport:  
['No' 'Yes' 'No internet service']
```

```
Unique values in StreamingTV:  
['No' 'Yes' 'No internet service']
```

```
Unique values in StreamingMovies:  
['No' 'Yes' 'No internet service']
```

```
Unique values in Contract:  
['Month-to-month' 'One year' 'Two year']
```

```
Unique values in PaperlessBilling:  
['Yes' 'No']
```

```
Unique values in PaymentMethod:  
['Electronic check' 'Mailed check' 'Bank transfer (automatic)'  
'Credit card (automatic)']
```

```
Unique values in MonthlyCharges:  
[29.85 56.95 53.85 ... 63.1 44.2 78.7 ]
```

```
Unique values in TotalCharges:  
[ 29.85 1889.5 108.15 ... 346.45 306.6 6844.5 ]
```

```
Unique values in Churn:  
['No' 'Yes']
```

```
[15]: # Example: Standardize 'Yes'/'No' to lowercase  
df = df.apply(lambda col: col.map(lambda s: s.lower() if type(s) == str else s))
```

```
[17]: # Example: Standardize columns with similar values (e.g., 'Male'/'Female' to  
#       'male'/'female')  
if 'gender' in df.columns:  
    df['gender'] = df['gender'].str.lower()
```

```
[19]: # Display the first few rows after standardization  
print(df.head())
```

```
customerID  gender  SeniorCitizen Partner Dependents  tenure PhoneService \
0  7590-vhveg  female           0      yes          no        1         no
```

```

1 5575-gnvde male 0 no no 34 yes
2 3668-qpybk male 0 no no 2 yes
3 7795-cfocw male 0 no no 45 no
4 9237-hqitu female 0 no no 2 yes

    MultipleLines InternetService OnlineSecurity ... DeviceProtection \
0 no phone service dsl no ... no
1 no dsl yes ... yes
2 no dsl yes ... no
3 no phone service dsl yes ... yes
4 no fiber optic no ... no

TechSupport StreamingTV StreamingMovies Contract PaperlessBilling \
0 no no no month-to-month yes
1 no no no one year no
2 no no no month-to-month yes
3 yes no no one year no
4 no no no month-to-month yes

PaymentMethod MonthlyCharges TotalCharges Churn
0 electronic check 29.85 29.85 no
1 mailed check 56.95 1889.50 no
2 mailed check 53.85 108.15 yes
3 bank transfer (automatic) 42.30 1840.75 no
4 electronic check 70.70 151.65 yes

```

[5 rows x 21 columns]

```
[21]: # Save the cleaned data
df.to_csv('Telco-Customer-Churn-cleaned.csv', index=False)
```

0.1.6 6. Convert columns to the correct data types as needed.

```
[23]: # Convert columns to appropriate data types
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce') # ↴ Convert to numeric
df['Churn'] = df['Churn'].astype('category') # Convert to category
df['tenure'] = df['tenure'].astype('int') # Ensure tenure is integer
```

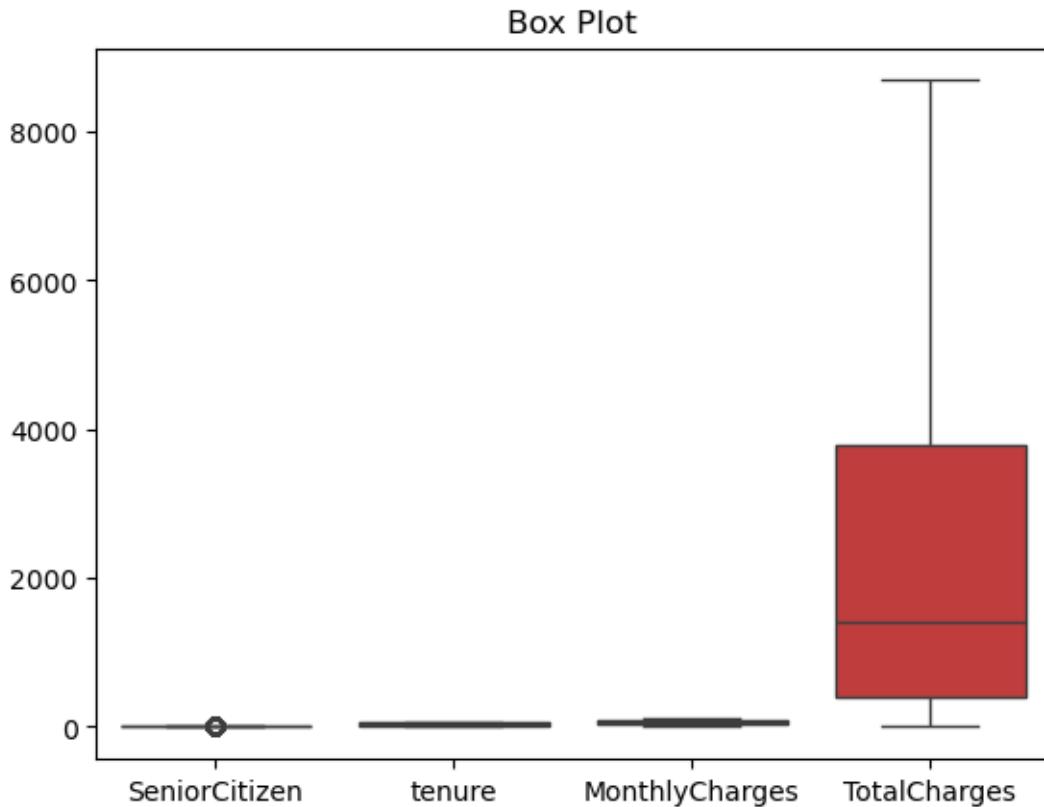
```
[25]: df.dtypes
```

customerID	object
gender	object
SeniorCitizen	int64
Partner	object
Dependents	object
tenure	int32

```
PhoneService          object
MultipleLines         object
InternetService       object
OnlineSecurity        object
OnlineBackup          object
DeviceProtection      object
TechSupport           object
StreamingTV          object
StreamingMovies       object
Contract              object
PaperlessBilling      object
PaymentMethod         object
MonthlyCharges        float64
TotalCharges          float64
Churn                 category
dtype: object
```

0.1.7 7. Identify and handle outliers in the data.

```
[27]: import seaborn as sns
# Box Plot
sns.boxplot(data=df)
plt.title("Box Plot")
plt.show()
```



0.1.8 8. Perform feature engineering, creating new features that may be relevant to predicting customer churn.

```
[29]: # Create age group feature (replace thresholds as needed)
df['age_group'] = np.where(df['SeniorCitizen'] == 1, 'Senior', np.
    where(df['tenure'] > 60, 'Middle-aged', 'Young'))

# Feature creation based on numerical features
# Tenure groups (replace thresholds as needed)
df['tenure_group'] = np.where(df['tenure'] <= 12, 'New Customer', 'Loyal
    Customer')

# Average monthly spend
df['avg_monthly_spend'] = df['MonthlyCharges'] / df['tenure']
```

0.1.9 9. Normalize or scale the data if necessary.

```
[31]: # Feature scaling or normalization (optional, consider data distribution)
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
numerical_features = ['tenure', 'MonthlyCharges', 'TotalCharges', ↴
    ↴'avg_monthly_spend'] # Adjust as needed
df[numerical_features] = scaler.fit_transform(df[numerical_features])
```

0.1.10 10. Split the dataset into training and testing sets for further analysis.

```
[33]: from sklearn.model_selection import train_test_split

# Assuming your data is in a DataFrame named 'df' and your target variable is ↴
    ↴ 'Churn' (replace if different)
X = df.drop('Churn', axis=1) # Features (all columns except 'Churn')
y = df['Churn'] # Target variable

# Split data into training and testing sets (common split is 80/20)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ↴
    ↴random_state=42)

print("Training data shapes:")
print(X_train.shape, y_train.shape)

print("Testing data shapes:")
print(X_test.shape, y_test.shape)
```

Training data shapes:

(5625, 23) (5625,)

Testing data shapes:

(1407, 23) (1407,)

```
[35]: #cleaned data is in a DataFrame named 'df'
df.to_csv('cleaned_data.csv', index=False) # Export without index
print("Cleaned data exported to cleaned_data.csv")
```

Cleaned data exported to cleaned_data.csv

[]: