

# CL1\_02

July 15, 2025

```
[ ]: """
NAME:Aher Swami Sandip
ROLL NO.01
COURSE: AI&DS
CLASS: BE
SUB:Computer Laboratory-I (Machine Learning)
"""
```

```
[ ]: """
PRACTICAL NO:02
A. Predict the price of the Uber ride from a given pickup point to the agreed ↴ drop-off
location. Perform following tasks:
1. Pre-process the dataset.
2. Identify outliers.
3. Check the correlation.
4. Implement linear regression and ridge, Lasso regression models.
5. Evaluate the models and compare their respective scores like R2, RMSE, etc. ↴
"""
"
```

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import math
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from sklearn.preprocessing import StandardScaler
```

```
[2]: data=pd.read_csv('uber.csv')
```

```
[4]: data.head()
```

```
[4]:   Unnamed: 0          key  fare_amount \
0    24238194  2015-05-07 19:52:06.0000003      7.5
1    27835199  2009-07-17 20:04:56.0000002      7.7
2    44984355  2009-08-24 21:45:00.00000061     12.9
```

```

3    25894730    2009-06-26 08:22:21.0000001      5.3
4    17610152    2014-08-28 17:47:00.000000188     16.0

          pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC        -73.999817       40.738354
1  2009-07-17 20:04:56 UTC        -73.994355       40.728225
2  2009-08-24 21:45:00 UTC        -74.005043       40.740770
3  2009-06-26 08:22:21 UTC        -73.976124       40.790844
4  2014-08-28 17:47:00 UTC        -73.925023       40.744085

  dropoff_longitude  dropoff_latitude  passenger_count
0        -73.999512        40.723217            1
1        -73.994710        40.750325            1
2        -73.962565        40.772647            1
3        -73.965316        40.803349            3
4        -73.973082        40.761247            5

```

[4]: `data.tail()`

```

[4]:           Unnamed: 0                 key  fare_amount
199995    42598914  2012-10-28 10:49:00.00000053      3.0 \
199996    16382965  2014-03-14 01:09:00.0000008       7.5
199997    27804658  2009-06-29 00:42:00.00000078      30.9
199998    20259894  2015-05-20 14:56:25.0000004      14.5
199999    11951496  2010-05-15 04:08:00.00000076      14.1

          pickup_datetime  pickup_longitude  pickup_latitude \
199995  2012-10-28 10:49:00 UTC        -73.987042       40.739367 \
199996  2014-03-14 01:09:00 UTC        -73.984722       40.736837
199997  2009-06-29 00:42:00 UTC        -73.986017       40.756487
199998  2015-05-20 14:56:25 UTC        -73.997124       40.725452
199999  2010-05-15 04:08:00 UTC        -73.984395       40.720077

  dropoff_longitude  dropoff_latitude  passenger_count
199995        -73.986525        40.740297            1
199996        -74.006672        40.739620            1
199997        -73.858957        40.692588            2
199998        -73.983215        40.695415            1
199999        -73.985508        40.768793            1

```

[5]: `data.dtypes`

```

[5]: Unnamed: 0      int64
key          object
fare_amount   float64
pickup_datetime  object
pickup_longitude float64

```

```
pickup_latitude      float64
dropoff_longitude   float64
dropoff_latitude    float64
passenger_count     int64
dtype: object
```

```
[6]: data.shape
```

```
[6]: (200000, 9)
```

```
[7]: data.describe()
```

```
[7]:      Unnamed: 0  fare_amount  pickup_longitude  pickup_latitude
count  2.000000e+05  200000.000000  200000.000000  200000.000000 \
mean   2.771250e+07   11.359955   -72.527638   39.935885
std    1.601382e+07   9.901776   11.437787   7.720539
min   1.000000e+00  -52.000000  -1340.648410  -74.015515
25%   1.382535e+07   6.000000  -73.992065   40.734796
50%   2.774550e+07   8.500000  -73.981823   40.752592
75%   4.155530e+07  12.500000  -73.967154   40.767158
max   5.542357e+07  499.000000   57.418457  1644.421482

      dropoff_longitude  dropoff_latitude  passenger_count
count      199999.000000  199999.000000  200000.000000
mean       -72.525292   39.923890   1.684535
std        13.117408   6.794829   1.385997
min      -3356.666300  -881.985513   0.000000
25%      -73.991407   40.733823   1.000000
50%      -73.980093   40.753042   1.000000
75%      -73.963658   40.768001   2.000000
max      1153.572603   872.697628  208.000000
```

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        200000 non-null  int64  
 1   key              200000 non-null  object  
 2   fare_amount      200000 non-null  float64 
 3   pickup_datetime  200000 non-null  object  
 4   pickup_longitude 200000 non-null  float64 
 5   pickup_latitude   200000 non-null  float64 
 6   dropoff_longitude 199999 non-null  float64 
 7   dropoff_latitude   199999 non-null  float64 
 8   passenger_count   200000 non-null  int64
```

```
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

```
[8]: data=data.drop(['Unnamed: 0','key','pickup_datetime'],axis=1)
```

```
[9]: data.isnull().sum()
```

```
[9]: fare_amount      0
pickup_longitude    0
pickup_latitude     0
dropoff_longitude   1
dropoff_latitude    1
passenger_count     0
dtype: int64
```

```
[10]: data=data.dropna()
```

```
[11]: data.isnull().sum()
```

```
[11]: fare_amount      0
pickup_longitude    0
pickup_latitude     0
dropoff_longitude   0
dropoff_latitude    0
passenger_count     0
dtype: int64
```

```
[12]: data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Index: 199999 entries, 0 to 199999
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   fare_amount      199999 non-null   float64
 1   pickup_longitude 199999 non-null   float64
 2   pickup_latitude  199999 non-null   float64
 3   dropoff_longitude 199999 non-null   float64
 4   dropoff_latitude  199999 non-null   float64
 5   passenger_count  199999 non-null   int64  
dtypes: float64(5), int64(1)
memory usage: 10.7 MB
```

```
[14]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 199999 entries, 0 to 199999
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   fare_amount      199999 non-null   float64
 1   pickup_longitude 199999 non-null   float64
 2   pickup_latitude  199999 non-null   float64
 3   dropoff_longitude 199999 non-null   float64
 4   dropoff_latitude  199999 non-null   float64
 5   passenger_count  199999 non-null   int64  
dtypes: float64(5), int64(1)
memory usage: 10.7 MB
```

```
---  
0    fare_amount      199999 non-null   float64  
1    pickup_longitude 199999 non-null   float64  
2    pickup_latitude   199999 non-null   float64  
3    dropoff_longitude 199999 non-null   float64  
4    dropoff_latitude   199999 non-null   float64  
5    passenger_count   199999 non-null   int64  
dtypes: float64(5), int64(1)  
memory usage: 10.7 MB
```

```
[13]: numeric_cols = data.select_dtypes(include=['int', 'float']).columns.tolist()  
print(numeric_cols)  
for col in numeric_cols:  
    sns.boxplot(data[numeric_cols])  
  
['fare_amount', 'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',  
'dropoff_latitude', 'passenger_count']
```



```
[14]: X = data.drop('fare_amount', axis=1)  
y = data['fare_amount']
```

```
[15]: Q1 = X.quantile(0.25)  
Q3 = X.quantile(0.75)
```

```
IQR = Q3 - Q1
outliers = ((X < (Q1 - 1.5 * IQR)) | (X > (Q3 + 1.5 * IQR))).any(axis=1)
X = X[~outliers]
y = y[~outliers]
```

```
[16]: correlation_matrix = data.corr()
for col in data.columns:
```

```
    print(correlation_matrix[col])
    print()
```

```
fare_amount      1.000000
pickup_longitude 0.010458
pickup_latitude   -0.008482
dropoff_longitude 0.008986
dropoff_latitude   -0.011014
passenger_count     0.010158
Name: fare_amount, dtype: float64
```

```
fare_amount      0.010458
pickup_longitude 1.000000
pickup_latitude   -0.816461
dropoff_longitude 0.833026
dropoff_latitude   -0.846324
passenger_count   -0.000415
Name: pickup_longitude, dtype: float64
```

```
fare_amount      -0.008482
pickup_longitude -0.816461
pickup_latitude    1.000000
dropoff_longitude -0.774787
dropoff_latitude    0.702367
passenger_count    -0.001559
Name: pickup_latitude, dtype: float64
```

```
fare_amount      0.008986
pickup_longitude 0.833026
pickup_latitude   -0.774787
dropoff_longitude 1.000000
dropoff_latitude   -0.917010
passenger_count    0.000033
Name: dropoff_longitude, dtype: float64
```

```
fare_amount      -0.011014
pickup_longitude -0.846324
pickup_latitude    0.702367
dropoff_longitude -0.917010
dropoff_latitude    1.000000
```


```

passenger_count      -0.000659
Name: dropoff_latitude, dtype: float64

fare_amount          0.010158
pickup_longitude    -0.000415
pickup_latitude     -0.001559
dropoff_longitude   0.000033
dropoff_latitude    -0.000659
passenger_count      1.000000
Name: passenger_count, dtype: float64

```

```
[17]: dataplot=sns.heatmap(data.corr(), annot=True)
plt.show()
```



```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
[19]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42)

[20]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

[21]: linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)

[21]: LinearRegression()

[22]: ridge_reg = Ridge(alpha=1.0)
ridge_reg.fit(X_train, y_train)

[22]: Ridge()

[23]: lasso_reg = Lasso(alpha=1.0)
lasso_reg.fit(X_train, y_train)

[23]: Lasso()

[24]: y_pred_linear = linear_reg.predict(X_test)
y_pred_ridge = ridge_reg.predict(X_test)
y_pred_lasso = lasso_reg.predict(X_test)

[25]: r2_linear = r2_score(y_test, y_pred_linear)
rmse_linear = np.sqrt(mean_squared_error(y_test, y_pred_linear))
mae_linear = mean_absolute_error(y_test, y_pred_linear)

r2_ridge = r2_score(y_test, y_pred_ridge)
rmse_ridge = np.sqrt(mean_squared_error(y_test, y_pred_ridge))
mae_ridge = mean_absolute_error(y_test, y_pred_ridge)

r2_lasso = r2_score(y_test, y_pred_lasso)
rmse_lasso = np.sqrt(mean_squared_error(y_test, y_pred_lasso))
mae_lasso = mean_absolute_error(y_test, y_pred_lasso)

[26]: print("Linear Regression - R2:", round(r2_linear,2), "RMSE:",round(rmse_linear), "MAE:", round(mae_linear))
print("Ridge Regression - R2:", round(r2_ridge), "RMSE:", round(rmse_ridge), "MAE:", round(mae_ridge))
print("Lasso Regression - R2:", round(r2_lasso), "RMSE:", round(rmse_lasso), "MAE:", round(mae_lasso))
```

Linear Regression - R2: 0.02 RMSE: 5 MAE: 3

Ridge Regression - R2: 0 RMSE: 5 MAE: 3

Lasso Regression - R2: 0 RMSE: 5 MAE: 3

```
[27]: #prediction of price  
#heatmap
```

```
[28]: import math  
print(abs(linear_reg.predict([[73.987042,40.739367,-73.986525,40.  
    ↪740297,1]]))[0]))  
print(abs(ridge_reg.predict([[73.987042,40.739367,-73.986525,40.740297,1]]))[0]))  
print(abs(lasso_reg.predict([[73.987042,40.739367,-73.986525,40.740297,1]]))[0]))
```

85.47108766462476

85.46891283895303

9.088837443709256