

Practical No: 1

Name: Thorave Avishkar Shrikrushna

Roll No: 65

Class: BE AI&DS

Title : To use PCA Algorithm for dimensionality reduction. You have a dataset that includes measurements for different variables on wine (alcohol, ash, magnesium, and so on).

Apply PCA algorithm & transform this data so that most variations in the measurements of the variables are captured by a small number of principal components so that it is easier to distinguish between red and white wine by inspecting these principal components.

Subject : Computer Laboratory 1 (Machine Learning) 417525

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
```

In [3]:

```
data=pd.read_csv('Wine.csv')
```

In [5]:

```
data.head()
```

Out[5]:

|   | Alcohol | Malic_Acid | Ash  | Ash_Alcanity | Magnesium | Total_Phenols | Flavanoids | Non |
|---|---------|------------|------|--------------|-----------|---------------|------------|-----|
| 0 | 14.23   | 1.71       | 2.43 | 15.6         | 127       | 2.80          | 3.06       |     |
| 1 | 13.20   | 1.78       | 2.14 | 11.2         | 100       | 2.65          | 2.76       |     |
| 2 | 13.16   | 2.36       | 2.67 | 18.6         | 101       | 2.80          | 3.24       |     |
| 3 | 14.37   | 1.95       | 2.50 | 16.8         | 113       | 3.85          | 3.49       |     |
| 4 | 13.24   | 2.59       | 2.87 | 21.0         | 118       | 2.80          | 2.69       |     |

◀ ━━━━ ▶

In [7]:

```
data.tail()
```

Out[7]:

|     | Alcohol | Malic_Acid | Ash  | Ash_Alcanity | Magnesium | Total_Phenols | Flavanoids | N |
|-----|---------|------------|------|--------------|-----------|---------------|------------|---|
| 173 | 13.71   | 5.65       | 2.45 | 20.5         | 95        | 1.68          | 0.61       |   |
| 174 | 13.40   | 3.91       | 2.48 | 23.0         | 102       | 1.80          | 0.75       |   |
| 175 | 13.27   | 4.28       | 2.26 | 20.0         | 120       | 1.59          | 0.69       |   |
| 176 | 13.17   | 2.59       | 2.37 | 20.0         | 120       | 1.65          | 0.68       |   |
| 177 | 14.13   | 4.10       | 2.74 | 24.5         | 96        | 2.05          | 0.76       |   |

◀ ━━━━ ▶

In [9]:

```
data.shape
```

Out[9]:

```
(178, 14)
```

In [11]:

```
data.describe()
```

Out[11]:

|              | Alcohol    | Malic_Acid | Ash        | Ash_Alcanity | Magnesium  | Total_Phenols | F1 |
|--------------|------------|------------|------------|--------------|------------|---------------|----|
| <b>count</b> | 178.000000 | 178.000000 | 178.000000 | 178.000000   | 178.000000 | 178.000000    | 17 |
| <b>mean</b>  | 13.000618  | 2.336348   | 2.366517   | 19.494944    | 99.741573  | 2.295112      |    |
| <b>std</b>   | 0.811827   | 1.117146   | 0.274344   | 3.339564     | 14.282484  | 0.625851      |    |
| <b>min</b>   | 11.030000  | 0.740000   | 1.360000   | 10.600000    | 70.000000  | 0.980000      |    |
| <b>25%</b>   | 12.362500  | 1.602500   | 2.210000   | 17.200000    | 88.000000  | 1.742500      |    |
| <b>50%</b>   | 13.050000  | 1.865000   | 2.360000   | 19.500000    | 98.000000  | 2.355000      |    |
| <b>75%</b>   | 13.677500  | 3.082500   | 2.557500   | 21.500000    | 107.000000 | 2.800000      |    |
| <b>max</b>   | 14.830000  | 5.800000   | 3.230000   | 30.000000    | 162.000000 | 3.880000      |    |

◀ ▶

In [13]: `data['Customer_Segment'].unique()`

Out[13]: `array([1, 2, 3], dtype=int64)`

In [15]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 178 entries, 0 to 177
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Alcohol          178 non-null    float64
 1   Malic_Acid       178 non-null    float64
 2   Ash               178 non-null    float64
 3   Ash_Alcanity     178 non-null    float64
 4   Magnesium         178 non-null    int64  
 5   Total_Phenols    178 non-null    float64
 6   Flavanoids        178 non-null    float64
 7   Nonflavanoid_Phenols  178 non-null    float64
 8   Proanthocyanins  178 non-null    float64
 9   Color_Intensity  178 non-null    float64
 10  Hue              178 non-null    float64
 11  OD280            178 non-null    float64
 12  Proline           178 non-null    int64  
 13  Customer_Segment 178 non-null    int64  
dtypes: float64(11), int64(3)
memory usage: 19.6 KB
```

In [17]: `data.isnull().sum()`

```
Out[17]: Alcohol          0  
         Malic_Acid        0  
         Ash              0  
         Ash_Alcanity     0  
         Magnesium         0  
         Total_Phenols     0  
         Flavanoids         0  
         Nonflavanoid_Phenols 0  
         Proanthocyanins    0  
         Color_Intensity    0  
         Hue               0  
         OD280             0  
         Proline            0  
         Customer_Segment    0  
         dtype: int64
```

```
In [19]: x=data.drop('Customer_Segment',axis=1)  
y=data['Customer_Segment']
```

```
In [21]: x
```

```
Out[21]:   Alcohol  Malic_Acid  Ash  Ash_Alcanity  Magnesium  Total_Phenols  Flavanoids  N  
0      14.23       1.71  2.43           15.6      127        2.80      3.06  
1      13.20       1.78  2.14           11.2      100        2.65      2.76  
2      13.16       2.36  2.67           18.6      101        2.80      3.24  
3      14.37       1.95  2.50           16.8      113        3.85      3.49  
4      13.24       2.59  2.87           21.0      118        2.80      2.69  
...  
173    13.71       5.65  2.45           20.5      95        1.68      0.61  
174    13.40       3.91  2.48           23.0      102        1.80      0.75  
175    13.27       4.28  2.26           20.0      120        1.59      0.69  
176    13.17       2.59  2.37           20.0      120        1.65      0.68  
177    14.13       4.10  2.74           24.5      96        2.05      0.76
```

178 rows × 13 columns

```
In [23]: x.shape
```

```
Out[23]: (178, 13)
```

```
In [25]: x_standardized = (x - x.mean()) / x.std()
```

```
In [27]: y
```

```
Out[27]: 0      1
          1      1
          2      1
          3      1
          4      1
          ..
         173     3
         174     3
         175     3
         176     3
         177     3
Name: Customer_Segment, Length: 178, dtype: int64
```

```
In [37]: pca=PCA(n_components=3)
```

```
In [39]: x_pca=pca.fit_transform(x)
```

```
In [41]: x_pca.shape
```

```
Out[41]: (178, 3)
```

```
In [43]: pca_df = pd.DataFrame(x_pca, columns = ['pca_col1','pca_col2','pca_col3'])
```

```
In [45]: pca_df
```

```
Out[45]:
```

|     | pca_col1    | pca_col2  | pca_col3  |
|-----|-------------|-----------|-----------|
| 0   | 318.562979  | 21.492131 | -3.130735 |
| 1   | 303.097420  | -5.364718 | -6.822835 |
| 2   | 438.061133  | -6.537309 | 1.113223  |
| 3   | 733.240139  | 0.192729  | 0.917257  |
| 4   | -11.571428  | 18.489995 | 0.554422  |
| ... | ...         | ...       | ...       |
| 173 | -6.980211   | -4.541137 | 2.474707  |
| 174 | 3.131605    | 2.335191  | 4.309931  |
| 175 | 88.458074   | 18.776285 | 2.237577  |
| 176 | 93.456242   | 18.670819 | 1.788392  |
| 177 | -186.943190 | -0.213331 | 5.630510  |

178 rows × 3 columns

```
In [47]: pca.explained_variance_ratio_
```

```
Out[47]: array([9.98091230e-01, 1.73591562e-03, 9.49589576e-05])
```

Practical No: 2

Name : Thorave Avishkar Shrikrushna

Roll No : 65

Class: BE AI&DS

Title : Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

Perform following tasks: 1. Pre-process the dataset. 2. Identify outliers. 3. Check the correlation.

4. Implement linear regression and ridge, Lasso regression models. 5.

Evaluate the models and compare their respective scores like R2, RMSE, etc. Dataset link:

<https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>

Subject : Computer Laboratory 1 (Machine Learning) 417525

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, r2_score
```

In [3]:

```
# 1. Load and Pre-process the Dataset
df = pd.read_csv('uber.csv') # Change the name to the correct csv file
```

In [4]:

```
print("Initial Data Info:")
print(df.info())
print(df.head())
```

```

Initial Data Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        200000 non-null   int64  
 1   key               200000 non-null   object  
 2   fare_amount       200000 non-null   float64 
 3   pickup_datetime   200000 non-null   object  
 4   pickup_longitude  200000 non-null   float64 
 5   pickup_latitude   200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude  199999 non-null   float64 
 8   passenger_count   200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
None
      Unnamed: 0          key  fare_amount \
0    24238194  2015-05-07 19:52:06.0000003    7.5
1    27835199  2009-07-17 20:04:56.0000002    7.7
2    44984355  2009-08-24 21:45:00.00000061   12.9
3    25894730  2009-06-26 08:22:21.0000001    5.3
4    17610152  2014-08-28 17:47:00.000000188   16.0

      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC          -73.999817        40.738354
1  2009-07-17 20:04:56 UTC          -73.994355        40.728225
2  2009-08-24 21:45:00 UTC          -74.005043        40.740770
3  2009-06-26 08:22:21 UTC          -73.976124        40.790844
4  2014-08-28 17:47:00 UTC          -73.925023        40.744085

      dropoff_longitude  dropoff_latitude  passenger_count
0            -73.999512        40.723217             1
1            -73.994710        40.750325             1
2            -73.962565        40.772647             1
3            -73.965316        40.803349             3
4            -73.973082        40.761247             5

```

```
In [5]: # Drop rows with missing values
df = df.dropna()
```

```
In [6]: # Convert datetime columns
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

```
In [7]: # Feature Engineering: Extract hour, day, month
df['hour'] = df['pickup_datetime'].dt.hour
df['day'] = df['pickup_datetime'].dt.day
df['month'] = df['pickup_datetime'].dt.month
```

```
In [8]: # Remove unneeded columns
df = df.drop(columns=['key', 'pickup_datetime'])
```

```
In [9]: # Calculate distance (Haversine formula)
def haversine(lat1, lon1, lat2, lon2):
    R = 6371 # Earth radius in km
    phi1, phi2 = np.radians(lat1), np.radians(lat2)
    dphi = np.radians(lat2 - lat1)
    dlambd = np.radians(lon2 - lon1)
```

```

    a = np.sin(dphi/2)**2 + np.cos(phi1)*np.cos(phi2)*np.sin(dlambd/2)**2
    return 2 * R * np.arcsin(np.sqrt(a))

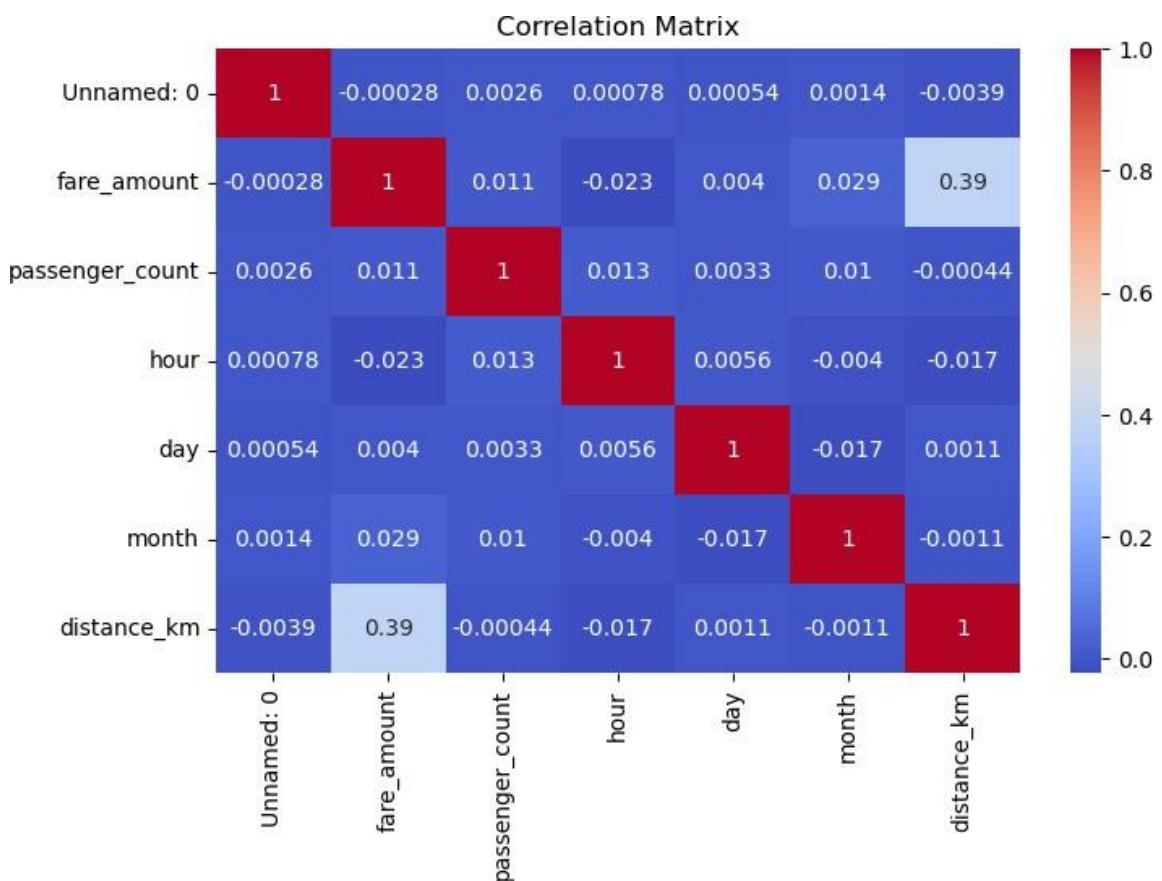
df['distance_km'] = haversine(df['pickup_latitude'], df['pickup_longitude'],
                               df['dropoff_latitude'], df['dropoff_longitude'])

```

In [10]: # Drop original lat/lon columns  
`df = df.drop(columns=['pickup_latitude', 'pickup_longitude', 'dropoff_latitude', 'dropoff_longitude'])`

In [11]: # 2. Identify Outliers (using z-score on fare\_amount and distance\_km)  
`from scipy.stats import zscore  
df = df[(np.abs(zscore(df[['fare_amount', 'distance_km']])) < 3).all(axis=1)]`

In [12]: # 3. Check Correlation  
`plt.figure(figsize=(8, 5))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()`



In [13]: # 4. Regression Models  
`X = df.drop(columns=['fare_amount'])  
y = df['fare_amount']`

In [14]: # One-hot encode categorical columns (if any)  
`X = pd.get_dummies(X, drop_first=True)`

In [15]: # Split data  
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_`

In [16]: # Linear Regression  
`lr = LinearRegression()`

```
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)
```

```
In [18]: # Ridge Regression
ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)
y_pred_ridge = ridge.predict(X_test)
```

```
In [19]: # Lasso Regression
lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
y_pred_lasso = lasso.predict(X_test)
```

```
In [21]: # 5. Model Evaluation
def print_scores(model_name, y_true, y_pred):
    print(f"\n{model_name}")
    print(f"R2 Score: {r2_score(y_true, y_pred):.4f}")
    print(f"RMSE: {np.sqrt(mean_squared_error(y_true, y_pred)):.4f}")

In [22]: print_scores("Linear Regression", y_test, y_pred_lr)
print_scores("Ridge Regression", y_test, y_pred_ridge)
print_scores("Lasso Regression", y_test, y_pred_lasso)
```

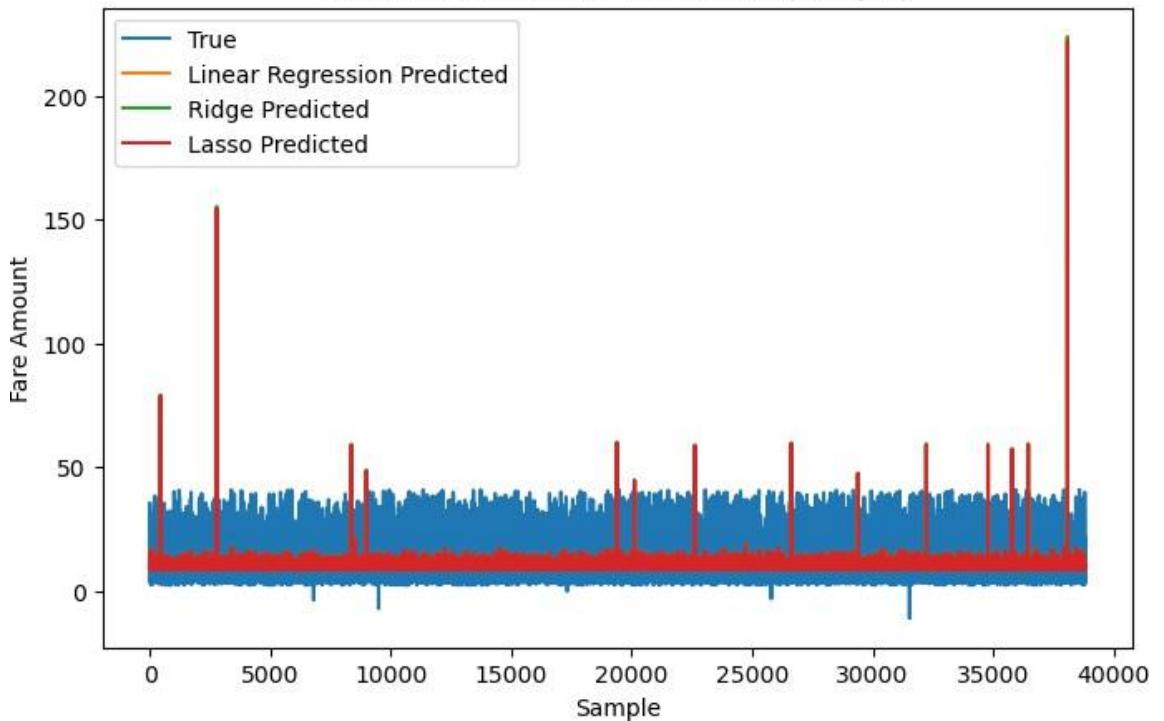
Linear Regression  
R2 Score: 0.1917  
RMSE: 5.8158

Ridge Regression  
R2 Score: 0.1917  
RMSE: 5.8158

Lasso Regression  
R2 Score: 0.1909  
RMSE: 5.8187

```
In [25]: # Optional: Compare visually
plt.figure(figsize=(8,5))
plt.plot(y_test.values, label='True')
plt.plot(y_pred_lr, label='Linear Regression Predicted')
plt.plot(y_pred_ridge, label='Ridge Predicted')
plt.plot(y_pred_lasso, label='Lasso Predicted')
plt.title('Model Predictions vs True Values (Sample)')
plt.xlabel('Sample')
plt.ylabel('Fare Amount')
plt.legend()
plt.show()
```

Model Predictions vs True Values (Sample)





Practical No: 4

Name: Thorave Avishkar Shrikrushna

Roll No: 65

Class: BE AI&DS

Title : Implement K-Means clustering on Iris.csv dataset. Determine the number of clusters using the elbow method.

Subject : Computer Laboratory 1 (Machine Learning) 417525

```
In [27]: import pandas as pd # Pandas (version : 1.1.5)
import numpy as np # Numpy (version : 1.19.2)
import matplotlib.pyplot as plt # Matplotlib (version : 3.3.2)
from sklearn.cluster import KMeans # Scikit Learn (version : 0.23.2)
import seaborn as sns # Seaborn (version : 0.11.1)
sns.set()
```

```
In [35]: data = pd.read_csv('iris.csv')
print(data)
```

|     | sepal_length | sepal_width | petal_length | petal_width | species   |
|-----|--------------|-------------|--------------|-------------|-----------|
| 0   | 5.1          | 3.5         | 1.4          | 0.2         | setosa    |
| 1   | 4.9          | 3.0         | 1.4          | 0.2         | setosa    |
| 2   | 4.7          | 3.2         | 1.3          | 0.2         | setosa    |
| 3   | 4.6          | 3.1         | 1.5          | 0.2         | setosa    |
| 4   | 5.0          | 3.6         | 1.4          | 0.2         | setosa    |
| ..  | ...          | ...         | ...          | ...         | ...       |
| 145 | 6.7          | 3.0         | 5.2          | 2.3         | virginica |
| 146 | 6.3          | 2.5         | 5.0          | 1.9         | virginica |
| 147 | 6.5          | 3.0         | 5.2          | 2.0         | virginica |
| 148 | 6.2          | 3.4         | 5.4          | 2.3         | virginica |
| 149 | 5.9          | 3.0         | 5.1          | 1.8         | virginica |

[150 rows x 5 columns]

```
In [31]: data.head()
```

```
Out[31]:
```

|   | sepal_length | sepal_width | petal_length | petal_width | species |
|---|--------------|-------------|--------------|-------------|---------|
| 0 | 5.1          | 3.5         | 1.4          | 0.2         | setosa  |
| 1 | 4.9          | 3.0         | 1.4          | 0.2         | setosa  |
| 2 | 4.7          | 3.2         | 1.3          | 0.2         | setosa  |
| 3 | 4.6          | 3.1         | 1.5          | 0.2         | setosa  |
| 4 | 5.0          | 3.6         | 1.4          | 0.2         | setosa  |

```
In [37]: data.tail()
```

```
Out[37]:
```

|     | sepal_length | sepal_width | petal_length | petal_width | species   |
|-----|--------------|-------------|--------------|-------------|-----------|
| 145 | 6.7          | 3.0         | 5.2          | 2.3         | virginica |
| 146 | 6.3          | 2.5         | 5.0          | 1.9         | virginica |
| 147 | 6.5          | 3.0         | 5.2          | 2.0         | virginica |
| 148 | 6.2          | 3.4         | 5.4          | 2.3         | virginica |
| 149 | 5.9          | 3.0         | 5.1          | 1.8         | virginica |

```
In [39]: len(data)
```

```
Out[39]: 150
```

```
In [41]: data.shape
```

```
Out[41]: (150, 5)
```

```
In [43]: data.columns
```

```
Out[43]: Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
       'species'],
       dtype='object')
```

```
In [47]: for i,col in enumerate(data.columns):
    print(f'Column number {1+i} is {col}')
```

```
Column number 1 is sepal_length
Column number 2 is sepal_width
Column number 3 is petal_length
Column number 4 is petal_width
Column number 5 is species
```

```
In [49]: data.dtypes
```

```
Out[49]: sepal_length    float64
          sepal_width     float64
          petal_length    float64
          petal_width     float64
          species         object
          dtype: object
```

```
In [51]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
 ----  -- 
 0   sepal_length    150 non-null    float64 
 1   sepal_width     150 non-null    float64 
 2   petal_length    150 non-null    float64 
 3   petal_width     150 non-null    float64 
 4   species         150 non-null    object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [53]: data.describe()
```

Out[53]:

|              | sepal_length | sepal_width | petal_length | petal_width |
|--------------|--------------|-------------|--------------|-------------|
| <b>count</b> | 150.000000   | 150.000000  | 150.000000   | 150.000000  |
| <b>mean</b>  | 5.843333     | 3.054000    | 3.758667     | 1.198667    |
| <b>std</b>   | 0.828066     | 0.433594    | 1.764420     | 0.763161    |
| <b>min</b>   | 4.300000     | 2.000000    | 1.000000     | 0.100000    |
| <b>25%</b>   | 5.100000     | 2.800000    | 1.600000     | 0.300000    |
| <b>50%</b>   | 5.800000     | 3.000000    | 4.350000     | 1.300000    |
| <b>75%</b>   | 6.400000     | 3.300000    | 5.100000     | 1.800000    |
| <b>max</b>   | 7.900000     | 4.400000    | 6.900000     | 2.500000    |

In [55]: `#Checking data for missing values using isnull()  
data.isnull()`

Out[55]:

|            | sepal_length | sepal_width | petal_length | petal_width | species |
|------------|--------------|-------------|--------------|-------------|---------|
| <b>0</b>   | False        | False       | False        | False       | False   |
| <b>1</b>   | False        | False       | False        | False       | False   |
| <b>2</b>   | False        | False       | False        | False       | False   |
| <b>3</b>   | False        | False       | False        | False       | False   |
| <b>4</b>   | False        | False       | False        | False       | False   |
| ...        | ...          | ...         | ...          | ...         | ...     |
| <b>145</b> | False        | False       | False        | False       | False   |
| <b>146</b> | False        | False       | False        | False       | False   |
| <b>147</b> | False        | False       | False        | False       | False   |
| <b>148</b> | False        | False       | False        | False       | False   |
| <b>149</b> | False        | False       | False        | False       | False   |

150 rows × 5 columns

In [69]: `print(data.columns.tolist())`

`['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species']`

In [71]: `data = pd.read_csv('iris.csv',header=0)`

In [77]: `data = pd.read_csv('iris.csv',header=None)  
data.columns=["Id","sepal_length","sepal_width","petal_length","petal_width"]`

In [79]: `print(data.columns.tolist())`

`['Id', 'sepal_length', 'sepal_width', 'petal_length', 'petal_width']`

In [83]: `#Checking summary of missing values`

```
data.isnull().sum()
```

```
Out[83]: Id          0  
sepal_length  0  
sepal_width   0  
petal_length   0  
petal_width   0  
dtype: int64
```

```
In [87]: #Deleting 'customer_id' column using drop().  
data.drop('Id', axis=1, inplace=True)  
data.head()
```

```
Out[87]:    sepal_length  sepal_width  petal_length  petal_width  
0      sepal_width    petal_length    petal_width       species  
1            3.5          1.4          0.2        setosa  
2            3.0          1.4          0.2        setosa  
3            3.2          1.3          0.2        setosa  
4            3.1          1.5          0.2        setosa
```

```
In [89]: data.isna().sum()
```

```
Out[89]: sepal_length  0  
sepal_width   0  
petal_length   0  
petal_width   0  
dtype: int64
```

```
In [91]: data.head()
```

```
Out[91]:    sepal_length  sepal_width  petal_length  petal_width  
0      sepal_width    petal_length    petal_width       species  
1            3.5          1.4          0.2        setosa  
2            3.0          1.4          0.2        setosa  
3            3.2          1.3          0.2        setosa  
4            3.1          1.5          0.2        setosa
```

```
In [111... data = pd.read_csv('iris.csv',header=None)  
data.columns=["sepal_length","sepal_width","petal_length","petal_width","Species"]
```

```
In [116... data.head()  
print(data.columns)
```

```
Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',  
       'Species'],  
      dtype='object')
```

```
In [118... data['Species'].value_counts()
```

```
Out[118... Species
setosa      50
versicolor   50
virginica    50
species      1
Name: count, dtype: int64
```

```
In [120... #Target Data
target_data = data.iloc[:,4]
target_data.head()
```

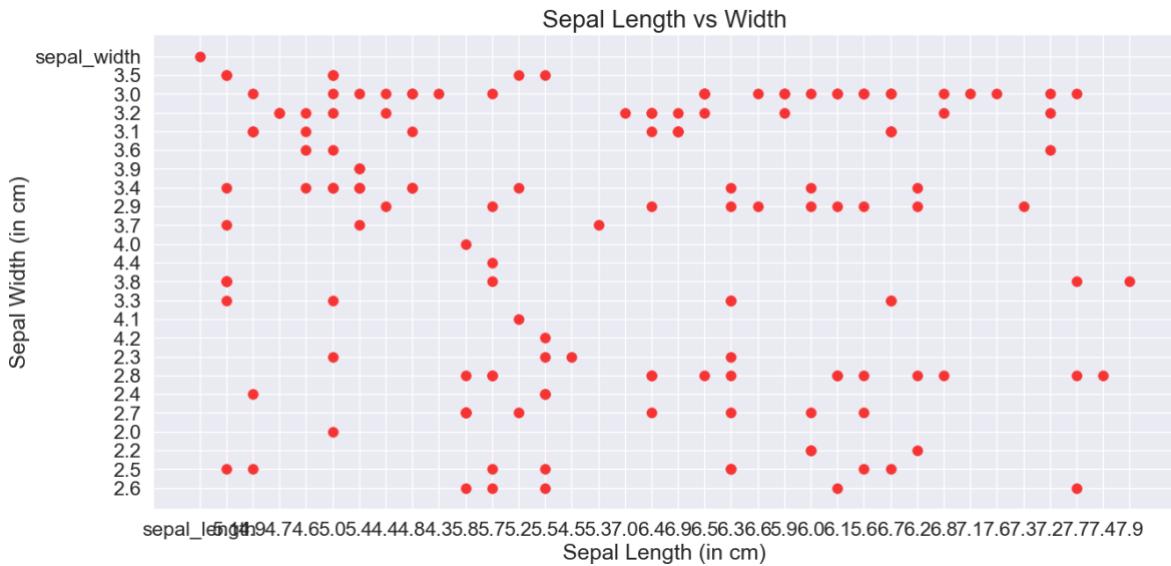
```
Out[120... 0    species
1    setosa
2    setosa
3    setosa
4    setosa
Name: Species, dtype: object
```

```
In [122... #Training data
clustering_data = data.iloc[:,[0,1,2,3]]
clustering_data.head()
```

|          | sepal_length | sepal_width | petal_length | petal_width |
|----------|--------------|-------------|--------------|-------------|
| <b>0</b> | sepal_length | sepal_width | petal_length | petal_width |
| <b>1</b> | 5.1          | 3.5         | 1.4          | 0.2         |
| <b>2</b> | 4.9          | 3.0         | 1.4          | 0.2         |
| <b>3</b> | 4.7          | 3.2         | 1.3          | 0.2         |
| <b>4</b> | 4.6          | 3.1         | 1.5          | 0.2         |

```
In [154... data.columns=["SepalLengthCm","SepalWidthCm","PetalLengthCm","PetalWidthCm","Spe
```

```
In [156... fig, ax = plt.subplots(figsize=(15,7))
sns.set(font_scale=1.5)
ax = sns.scatterplot(x=data['SepalLengthCm'],y=data['SepalWidthCm'], s=70, color
ax.set_ylabel('Sepal Width (in cm)')
ax.set_xlabel('Sepal Length (in cm)')
plt.title('Sepal Length vs Width', fontsize = 20)
plt.show()
```



In [188]:

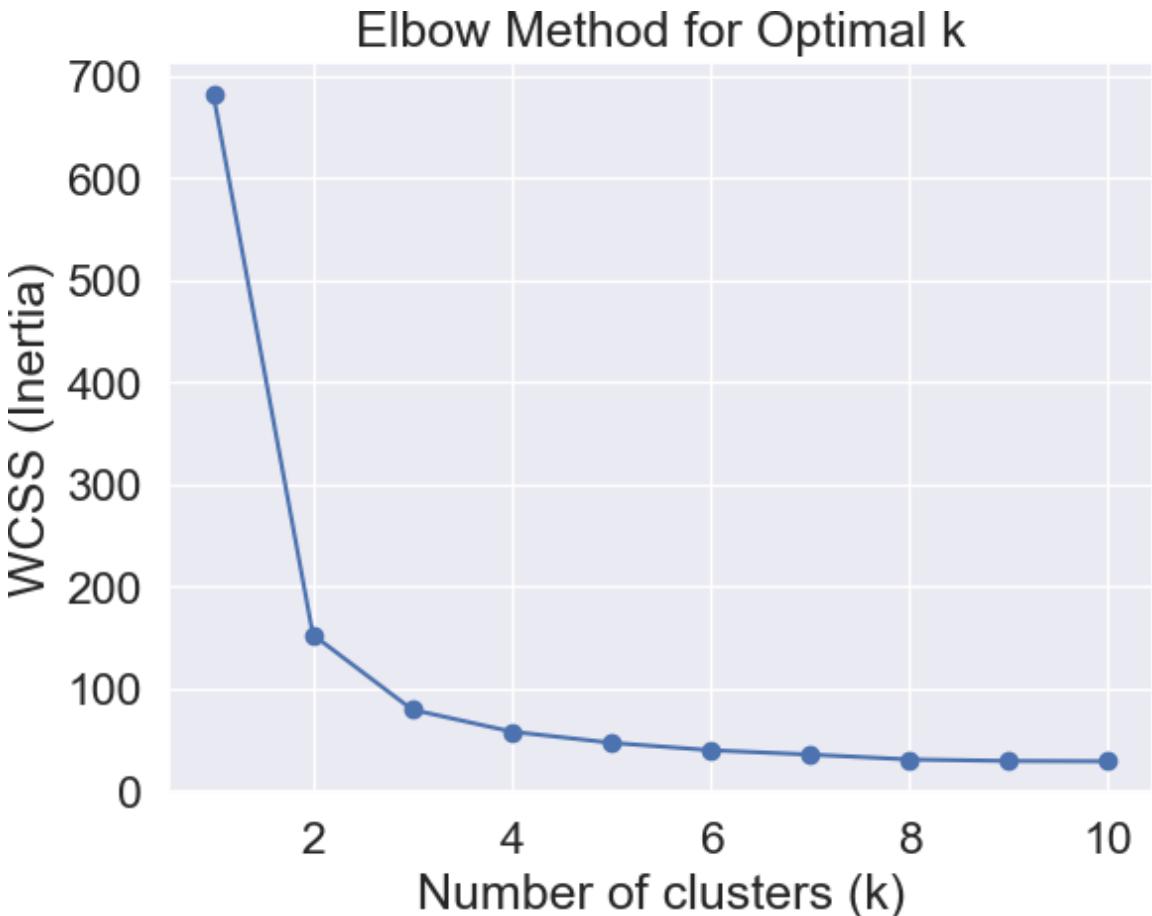
```
#The Elbow Method
import pandas as pd
import numpy as np
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Step 1: Load the dataset
clustering_data = pd.read_csv('iris.csv')

# Step 2: Select only numeric columns (exclude non-numeric like 'species' if present)
# If the dataset has columns: sepal_length, sepal_width, petal_length, petal_width
features = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
X = clustering_data[features]

# Step 3: Elbow method to find optimal number of clusters
wcss = []
for i in range(1, 11):
    km = KMeans(n_clusters=i, random_state=42)
    km.fit(X)
    wcss.append(km.inertia_)

# Step 4: Plot the elbow curve
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of clusters (k)')
plt.ylabel('WCSS (Inertia)')
plt.grid(True)
plt.show()
```



```
In [194...]: # clustering
from sklearn.cluster import KMeans

# Select only the numeric columns (exclude 'species')
features = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
X = clustering_data[features]

# Create and fit the model
kms = KMeans(n_clusters=3, init='k-means++', n_init='auto', random_state=42)
kms.fit(X)

# Create a copy and add cluster predictions
clusters = clustering_data.copy()
clusters['Cluster_Prediction'] = kms.predict(X)

# Show the result
clusters.head()
```

```
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
warnings.warn(
```

Out[194...]

|   | sep al_length | sepal_width | petal_length | petal_width | species | Cluster_Prediction |
|---|---------------|-------------|--------------|-------------|---------|--------------------|
| 0 | 5.1           | 3.5         | 1.4          | 0.2         | setosa  | 1                  |
| 1 | 4.9           | 3.0         | 1.4          | 0.2         | setosa  | 1                  |
| 2 | 4.7           | 3.2         | 1.3          | 0.2         | setosa  | 1                  |
| 3 | 4.6           | 3.1         | 1.5          | 0.2         | setosa  | 1                  |
| 4 | 5.0           | 3.6         | 1.4          | 0.2         | setosa  | 1                  |

In [196...]

```
kms.cluster_centers_
```

Out[196...]

```
array([[6.85384615, 3.07692308, 5.71538462, 2.05384615],
       [5.006      , 3.418      , 1.464      , 0.244      ],
       [5.88360656, 2.74098361, 4.38852459, 1.43442623]])
```

In [214...]

```
print(clusters.columns.tolist())
```

```
['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species', 'cluster_prediction']
```

In [238...]

```
from sklearn.cluster import KMeans

# Assume X is your input features, e.g., SepalLengthCm and SepalWidthCm
kms = KMeans(n_clusters=3, random_state=0)
clusters = X.copy() # Copy your input DataFrame

# Add predicted cluster labels
clusters['Cluster_Prediction'] = kms.fit_predict(X)
```

```
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
```

```
warnings.warn(
```

In [240...]

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

# Load dataset (example: Iris dataset)
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df.columns = ['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']

# Select features for clustering
X = df[['SepalLengthCm', 'SepalWidthCm']]

# Fit KMeans
kms = KMeans(n_clusters=3, random_state=42)
df['Cluster_Prediction'] = kms.fit_predict(X)

# Plotting
fig, ax = plt.subplots(figsize=(15, 7))

# Cluster 0
plt.scatter(
```

```

        df[df['Cluster_Prediction'] == 0]['SepalLengthCm'],
        df[df['Cluster_Prediction'] == 0]['SepalWidthCm'],
        s=70, c='teal', edgecolor='black', label='Cluster 0'
    )

    # Cluster 1
    plt.scatter(
        df[df['Cluster_Prediction'] == 1]['SepalLengthCm'],
        df[df['Cluster_Prediction'] == 1]['SepalWidthCm'],
        s=70, c='lime', edgecolor='black', label='Cluster 1'
    )

    # Cluster 2
    plt.scatter(
        df[df['Cluster_Prediction'] == 2]['SepalLengthCm'],
        df[df['Cluster_Prediction'] == 2]['SepalWidthCm'],
        s=70, c='magenta', edgecolor='black', label='Cluster 2'
    )

    # Plot centroids
    plt.scatter(
        kms.cluster_centers_[:, 0], kms.cluster_centers_[:, 1],
        s=170, c='yellow', edgecolor='black', label='Centroids'
    )

    # Labels, Limits, etc.
    plt.title('KMeans Clustering (Sepal Length vs Sepal Width)', fontsize=18)
    plt.xlabel('Sepal Length (cm)')
    plt.ylabel('Sepal Width (cm)')
    plt.xlim(4, 8)
    plt.ylim(1.8, 4.5)
    plt.legend()
    plt.grid(True)
    plt.tight_layout()
    plt.show()

```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.

```
warnings.warn(
```



In [ ]:



Practical No: 5

Name : Thorave Avishkar Shrikrushna

Roll No: 65

Class : BE AI&DS

Title : Use different voting mechanism and Apply AdaBoost (Adaptive Boosting), GradientTree Boosting (GBM), XGBoost classification on Iris dataset and compare the performance of three models using different evaluation measures.

Subject : Computer Laboratory 1 (Machine Learning) 417525

```
In [1]: import pandas as pd  
from sklearn.datasets import load_digits  
digits = load_digits()
```

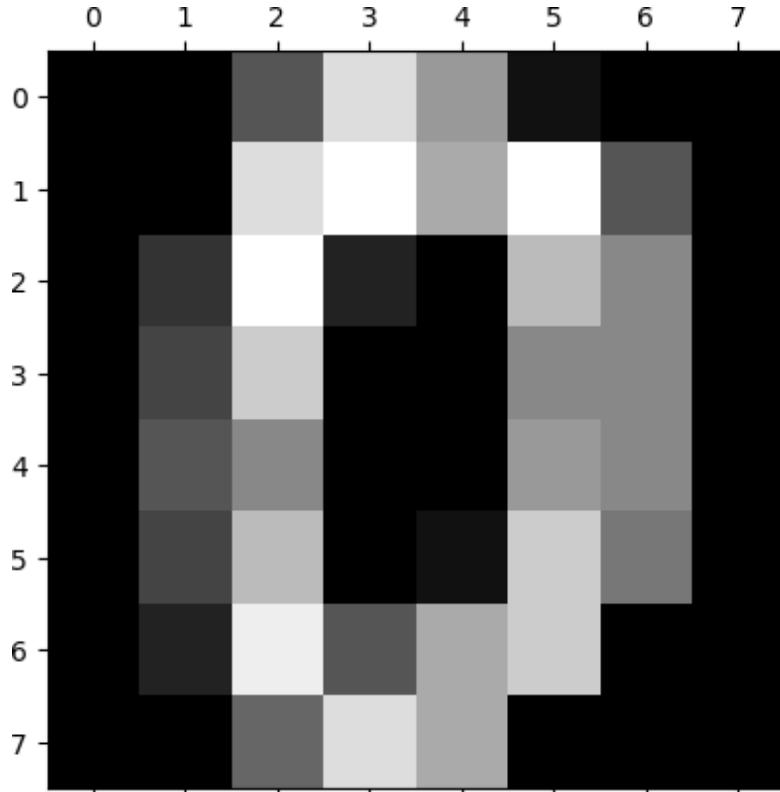
```
In [3]: dir(digits)
```

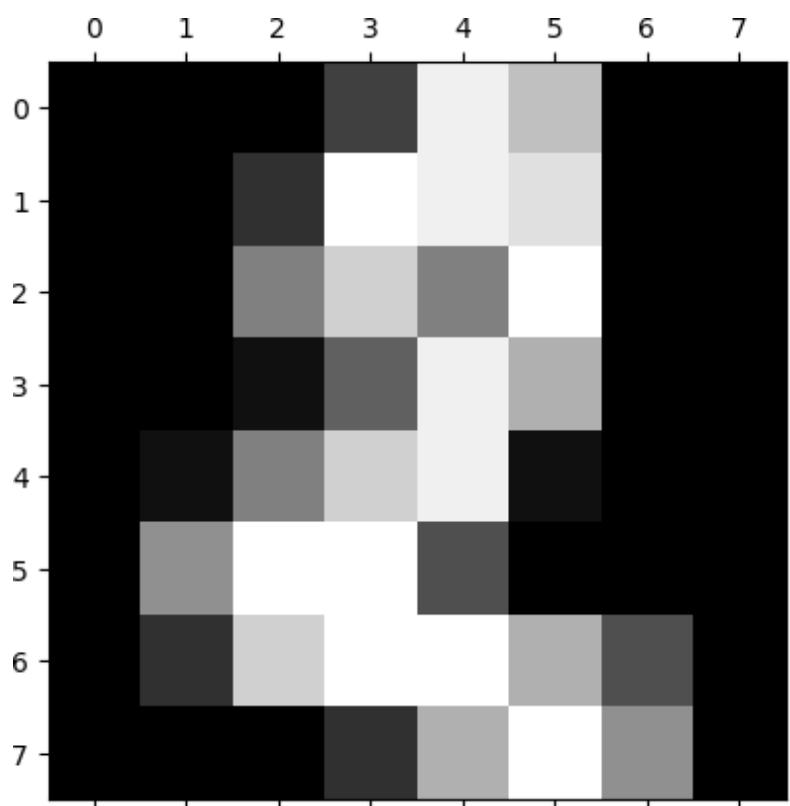
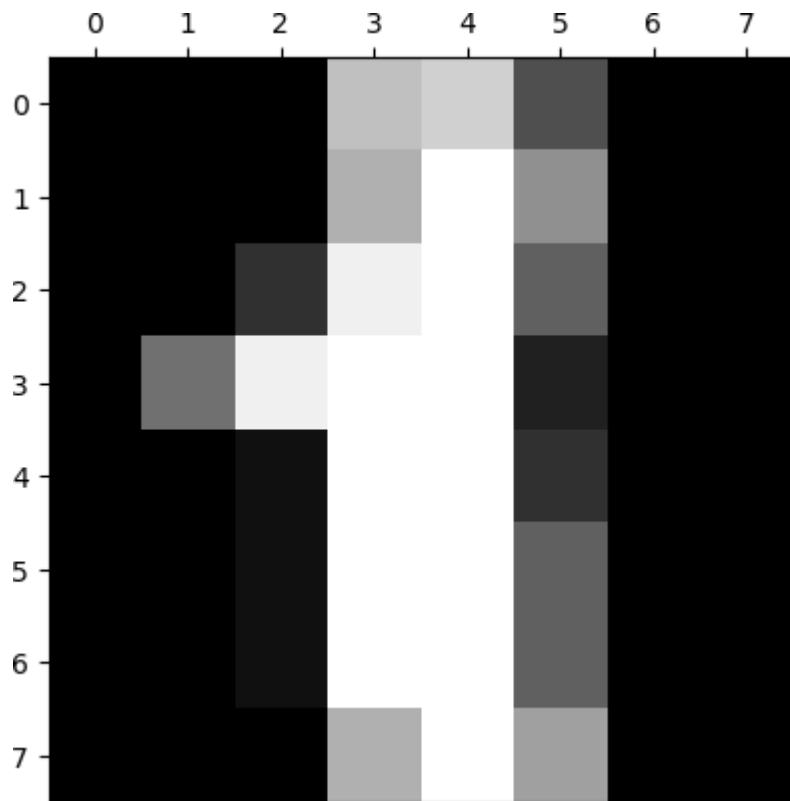
```
Out[3]: ['DESCR', 'data', 'feature_names', 'frame', 'images', 'target', 'target_names']
```

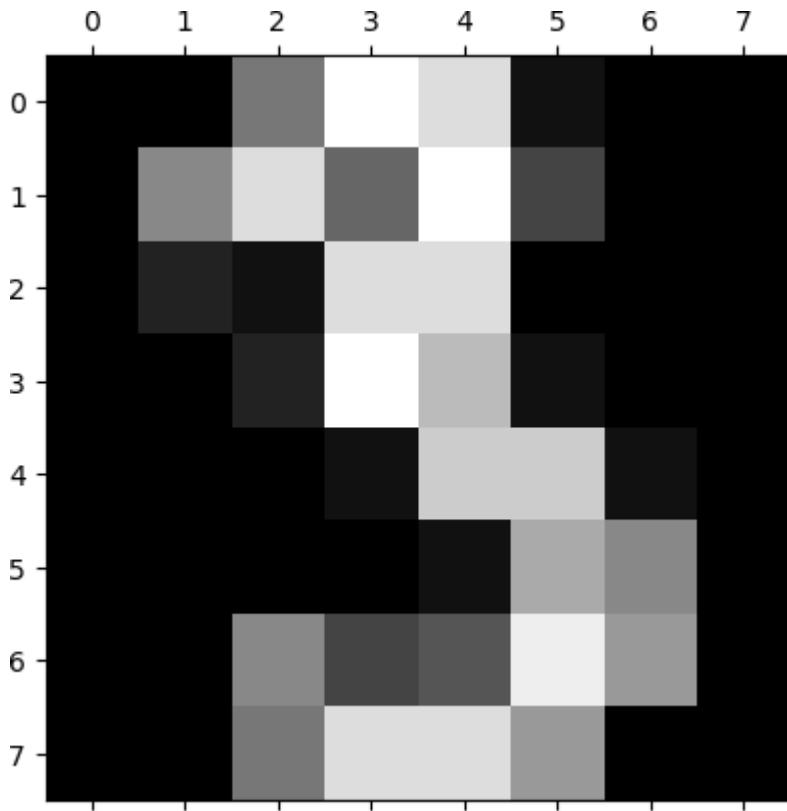
```
In [5]: #%matplotlib inline  
import matplotlib.pyplot as plt
```

```
In [9]: plt.gray()  
for i in range(4):  
    plt.matshow(digits.images[i])
```

<Figure size 640x480 with 0 Axes>







```
In [13]: df = pd.DataFrame(digits.data)
df.head()
```

```
Out[13]:      0   1   2   3   4   5   6   7   8   9   ...  54  55  56  57  58  59   6
  0  0.0  0.0  5.0  13.0  9.0  1.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  6.0  13.0  10.
  1  0.0  0.0  0.0  12.0  13.0  5.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  11.0  16.
  2  0.0  0.0  0.0  4.0   15.0  12.0  0.0  0.0  0.0  0.0  ...  5.0  0.0  0.0  0.0  0.0  3.0   11.
  3  0.0  0.0  7.0  15.0  13.0  1.0  0.0  0.0  0.0  8.0  ...  9.0  0.0  0.0  0.0  7.0  13.0  13.
  4  0.0  0.0  0.0  1.0  11.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  2.0  16.
```

5 rows × 64 columns



```
In [15]: df['target'] = digits.target
df[0:12]
```

Out[15]:

|    | 0   | 1   | 2    | 3    | 4    | 5    | 6    | 7   | 8   | 9   | ... | 55  | 56  | 57  | 58   | 59   | 60   |
|----|-----|-----|------|------|------|------|------|-----|-----|-----|-----|-----|-----|-----|------|------|------|
| 0  | 0.0 | 0.0 | 5.0  | 13.0 | 9.0  | 1.0  | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 6.0  | 13.0 | 10.0 |
| 1  | 0.0 | 0.0 | 0.0  | 12.0 | 13.0 | 5.0  | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0  | 11.0 | 16.0 |
| 2  | 0.0 | 0.0 | 0.0  | 4.0  | 15.0 | 12.0 | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0  | 3.0  | 11.0 |
| 3  | 0.0 | 0.0 | 7.0  | 15.0 | 13.0 | 1.0  | 0.0  | 0.0 | 0.0 | 8.0 | ... | 0.0 | 0.0 | 0.0 | 7.0  | 13.0 | 13.0 |
| 4  | 0.0 | 0.0 | 0.0  | 1.0  | 11.0 | 0.0  | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0  | 2.0  | 16.0 |
| 5  | 0.0 | 0.0 | 12.0 | 10.0 | 0.0  | 0.0  | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 9.0  | 16.0 | 16.0 |
| 6  | 0.0 | 0.0 | 0.0  | 12.0 | 13.0 | 0.0  | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 1.0  | 9.0  | 15.0 |
| 7  | 0.0 | 0.0 | 7.0  | 8.0  | 13.0 | 16.0 | 15.0 | 1.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 13.0 | 5.0  | 0.0  |
| 8  | 0.0 | 0.0 | 9.0  | 14.0 | 8.0  | 1.0  | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 11.0 | 16.0 | 15.0 |
| 9  | 0.0 | 0.0 | 11.0 | 12.0 | 0.0  | 0.0  | 0.0  | 0.0 | 0.0 | 2.0 | ... | 0.0 | 0.0 | 0.0 | 9.0  | 12.0 | 13.0 |
| 10 | 0.0 | 0.0 | 1.0  | 9.0  | 15.0 | 11.0 | 0.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 1.0  | 10.0 | 13.0 |
| 11 | 0.0 | 0.0 | 0.0  | 0.0  | 14.0 | 13.0 | 1.0  | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0  | 1.0  | 13.0 |

12 rows × 65 columns

In [21]:

```
#Train and the model and prediction
X = df.drop('target',axis='columns')
y = df.target
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2)
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=20)
model.fit(X_train, y_train)
RandomForestClassifier(n_estimators=20)
model.score(X_test, y_test)
```

Out[21]: 0.9722222222222222

In [23]:

```
y_predicted = model.predict(X_test)
```

In [25]:

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_predicted)
cm
```

Out[25]:

```
array([[34,  0,  0,  0,  1,  0,  0,  0,  0,  0],
       [ 0, 40,  0,  0,  0,  0,  0,  0,  0,  0],
       [ 0,  0, 34,  0,  0,  0,  0,  0,  1,  0],
       [ 0,  0,  1, 32,  0,  0,  0,  0,  0,  0],
       [ 0,  0,  0,  0, 33,  0,  0,  1,  0,  1],
       [ 0,  0,  0,  0,  0, 27,  0,  0,  0,  1],
       [ 1,  1,  0,  0,  0,  0, 38,  0,  0,  0],
       [ 0,  0,  0,  1,  0,  0,  0, 39,  0,  0],
       [ 0,  1,  0,  0,  0,  0,  0,  0, 29,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0,  0, 44]], dtype=int64)
```

In [27]:

```
%matplotlib inline
import matplotlib.pyplot as plt
```

```
import seaborn as sn
plt.figure(figsize=(10,7))
sn.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[27]: Text(95.72222222222221, 0.5, 'Truth')